

Stable Answers, Unfinished Reasoning: The Stability–Terminality Gap in Probe-Based Early Exit

Anonymous ACL submission

Abstract

A natural and increasingly studied way to cut reasoning-model inference cost is to probe the partial trajectory for its current answer and stop once probes *agree*. We show that repeated agreement is not a safe stopping criterion on its own, because it measures the wrong object: under a fixed probing procedure, agreement establishes that the current answer *persists*, not that the reasoning has *terminated*—a *stability–terminality gap*. The gap has a concrete source. A model that has not finished reasoning is still *forced* by the probe to emit an answer, and the cheapest thing to emit is what it emitted last time; this elicited placeholder can repeat for many probes and look settled, yet the trajectory goes on to correct it. We show this directly: re-probing the *same* frozen prefix with a differently worded query changes the extracted answer 57% of the time early in a trajectory (against 11% near the end), hand-labelling 134 wrong stops finds three in five commit to an answer the model had not converged on, and continued reasoning recovers a correct answer far more often than a stop banks a wrong one (up to $\sim 45\times$ at an aggressive latest-probe stop, $\sim 2\times$ at the largest windows, where the rule has nearly stopped firing). Widening the agreement window therefore does not make the signal better at detecting terminality; it only lowers the loss by firing later and less often. A *preregistered* sweep of 3,520 consensus rules (window size, share threshold, and operational knobs; frozen trajectories from two models and three math benchmarks; acceptance gates fixed in advance) shows no setting of these knobs closes the gap: **no rule clears any gate**. As a positive control that early exit is not itself impossible, a boundary-confidence method (DEER) swept through the identical pipeline **clears all three gates** and holds on a held-out split and two unseen models. Agreement fails not because it is insufficiently strict, but because it repeatedly measures the wrong object. We release the protocol, the frozen trajectories, and the

full sweep.

1 Introduction

Reasoning language models (DeepSeek-AI, 2025; Qwen Team, 2025; OpenAI, 2024) can spend thousands of tokens on a single answer. *Early exit*—halting a chain of thought once its reasoning is done—could cut that cost (Chen et al., 2024; Fu et al., 2024), but it requires knowing *when the reasoning has finished*, which the trajectory does not announce. A natural and increasingly studied proxy is *intermediate answer consensus*: at regular intervals a partial trajectory is probed for its current answer (e.g., by appending a short “Final Answer:” suffix), and generation is stopped once recent probes *agree* (Fu et al., 2024). The premise is that if the model repeatedly says x , it has probably finished reasoning about x ; recent work reports answers that “converge” well before a trajectory ends and stops on that convergence at little *aggregate* accuracy cost (Liu and Wang, 2025).

We show that this proxy measures the wrong object. Under a fixed probing procedure, repeated agreement establishes only that the current answer *persists*, not that the reasoning has *terminated*, and the two come apart—a *stability–terminality gap*.¹ This reconciles the aggregate savings above: those savings are real, but they ride on a *directional* accuracy cost that safety-blind, probe-cost-free evaluation averages away.

The gap has a concrete source. A model that has not finished reasoning is nonetheless *forced* by the probe to write an answer, and the cheapest thing to write is what it wrote last time; this elicited placeholder can repeat across many probes and look settled while the underlying trajectory goes on to correct it. Three direct measurements bear this out

¹“False consensus” here denotes non-terminal intermediate probe agreement, unrelated to the social-psychology term (Ross et al., 1977); we study answer-level, not token/layer-level, early exit (Graves, 2016; Schuster et al., 2022).

(Section 5): re-probing the *same* frozen prefix with a differently worded query returns a different answer 57% of the time in the first tenth of a trajectory against 11% near the end, so early “answers” are largely artifacts of being asked; hand-labelling 134 wrong stops finds three in five commit to an answer the model had not converged on, against one in five it had genuinely settled; and continued reasoning recovers a correct answer far more often than a stop banks a wrong one—up to $\sim 45\times$ at an aggressive latest-probe stop. Widening the agreement window does not close the gap: it lowers the loss only by firing later and less often, so at the largest windows the ratio falls to $\sim 2\times$ but the rule fires on a fifth as many problems and saves almost nothing.

To rule out that this is a badly chosen threshold or a single implementation, we search the signal exhaustively under commitment. On frozen trajectories from two models and three benchmarks we run a *preregistered* sweep of the windowed consensus family we study—window size $\times$ share threshold plus operational knobs, 3,520 rules—with problem-grouped splits and acceptance gates fixed in advance (Sections 3–4). No setting of these knobs closes the gap: **not one of the 3,520 rules is both safe and saving—none clears any gate**. The frontier reproduces on a held-out test split ($r=0.98$) and on unseen models at $4\times$ the scale and a different architecture, and a shipped operating point (CertaIndex (CoT)) loses 56–70 pp on the same trajectories.

That the gap cannot be tuned away leaves one alternative explanation—that early exit is simply impossible at these budgets. It is not. As a positive control we sweep one non-consensus method—DEER, which reads the model’s confidence at a reasoning boundary—through the identical pipeline and gates: it clears all three, losing as little as 0.3 pp while saving 28%, and does so on the unseen models too. DEER is a control, not a proposal; it establishes that early exit is achievable, so the consensus family’s failure is specific to *what agreement measures*. Figure 1 is the argument in one picture. Our claim concerns agreement used *alone* as a stopping criterion; combined with a signal that reflects the model’s own estimate it may still be useful, and we return to this in Section 5.

Contribution.

- **We identify a stability–terminality gap in probe-based early exit:** repeated intermediate agreement establishes that the current an-

swer *persists* under a fixed probing procedure, not that the reasoning has *terminated* (Section 5).

- **We explain what agreement actually measures,** through three direct signals—the wording sensitivity of the extracted answer, a hand-labelled error taxonomy of stopped-but-wrong cases ($\kappa=0.82$), and a trajectory-level recovery asymmetry—which together show that early agreement is a probe-elicited placeholder on an unfinished trajectory rather than a settled belief (Section 5).
- **We show the gap is not a threshold or implementation artifact:** a preregistered sweep of 3,520 rules on frozen trajectories closes no gate, the frontier reproduces on a held-out split and two unseen models, and a non-consensus positive control (DEER) swept under the identical protocol clears the same gates—ruling out that early exit is simply impossible (Sections 3–4). We release the protocol, trajectories, and sweep in full.

2 Related Work

Probe-based / consensus early exit. The methods we scrutinize periodically query a partial reasoning trajectory for its current answer and halt on agreement. Dynasor / CertaIndex (Fu et al., 2024) formalize this with a “certainty index” over recent probe answers and use it to terminate early when serving reasoning programs, one of several responses to the *overthinking* of reasoning models (Chen et al., 2024; Yang et al., 2025). Recent work pushes the same idea further: Liu and Wang (2025) report that a trajectory’s answer “converges” well before it ends and stop on that convergence at little *aggregate* accuracy cost, while learned prefix-feature stoppers reach the more cautious conclusion that no single aggressive policy is universally safe (Dong et al., 2026). We adopt the same probe mechanism (a short boxed-answer suffix) and notion of probe agreement, and ask the question these leave open: whether *any* rule in this family is both safe and saving, or whether the apparent aggregate savings hide a directional accuracy cost.

Where agreement gets its meaning. Self-consistency (Wang et al., 2023) samples a *diverse* set of reasoning paths and takes the majority answer, on the premise that a hard problem admits many routes to one correct answer. Agreement is




Figure 1: **The argument in one picture.** (a) Intermediate probe agreement is not terminal. On a real trajectory (MATH500 problem 68, DeepSeek-R1-Distill-Qwen-7B, dense 64-token probes) the model emits the same wrong answer, 37, for ten consecutive probes; a rule requiring three agreeing probes fires at generated token 256 and commits it. Left to run, the trajectory reaches the correct answer 46 and holds it for its final 47 probes. Ribbon colours: red the dominant wrong answer, pink other wrong answers, green correct. (b) We ask whether *any* rule in the searched family escapes this. A preregistered sweep of 3,520 rules—window size $W \times$ share threshold s , crossed with probe schedule, maturity, validity and certainty—is replayed on 18 frozen environments against three acceptance gates fixed before the sweep ran. None clears any gate, and the two ways to fail are symmetric: capping the accuracy drop at 1 pp permits 0.2% token saving, while demanding 10% saving costs 2.7 pp. (c) The failure is the *signal*, not early exit. Swept through the identical pipeline, gates and token accounting, a boundary-confidence method (DEER) clears all three gates while consensus clears none—and it does so on the held-out test split, at $4 \times$ scale and on a different architecture as well (§4.5).

evidence there because the paths are drawn independently: when routes that could have diverged do not, that is information about the answer, and the vote can outweigh an error on any single path. Early-stopping variants keep this structure and only decide how many samples to draw (Aggarwal et al., 2023; Li et al., 2024). CertIndex uses the same structure in its multi-path setting, measuring entropy over answers from independently sampled trajectories.

Probing a single chain of thought is a different object, and we call it *self-consensus* to keep the two apart. Its readings are not independent samples: every probe is conditioned on the same prefix, so agreement between them records that the prefix has not changed its answer, not that separate lines of reasoning have converged on one. This removes the protection self-consistency enjoys—agreement no longer supplies a vote that can outweigh an error—but on its own it does not settle the question that matters here: whether agreement, even so, tracks *terminality*. That the same-prefix answer is still moving is something we establish by direct measurement in Section 5—by re-wording the probe, labelling wrong stops, and tracking how often the trajectory later corrects itself—rather than

infer from non-independence alone. This within-trajectory answer evolution is also distinct from the prompted, post-hoc self-correction that Huang et al. (2024) find unreliable, and we find (Section 4) that a model’s mid-trajectory answer is far from final.

Alternative termination signals. Rather than probe agreement, a stop can read other signals: verifier- and process-reward-guided decoding (Lightman et al., 2024; Wang et al., 2024) score partial reasoning, and adaptive-computation methods learn a halting signal at the token or layer level (Graves, 2016; Schuster et al., 2022). Confidence-dynamics stops read how the model’s answer confidence evolves and halt when it settles—a forward-looking signal rather than backward probe agreement (Hosseini et al., 2026). DEER (Yang et al., 2025) stops on the model’s confidence in a trial answer at reasoning boundaries. We use DEER only as a *positive control* (Section 5): run through the identical pipeline and acceptance gates, it clears them where no consensus rule does, indicating that the failure is specific to the probe-agreement signal rather than to early exit itself. We do not claim this broader class of signals is categorically superior—our experiments isolate one signal, not a taxonomy

of methods. Test-time scaling trades compute for accuracy via repeated sampling and search (Snell et al., 2024; Brown et al., 2024); early exit is its dual, spending less without losing accuracy.

Preregistration. To keep an exhaustive rule search from silently overfitting to a favorable operating point, we fix the rule schema, splits, and acceptance gates before looking at held-out data (Nosek et al., 2025). This preregistration (Section 3) is what makes a *negative* result interpretable.

3 Method

To ask whether *any* probe-based rule can succeed, we need (i) an evaluation that a rule cannot game by changing the reasoning it observes, and (ii) a selection procedure that cannot silently overfit to a lucky operating point. We preregister both.

3.1 Probe-independent, frozen trajectories

For each problem we generate exactly *one* main trajectory in a single request (caps: 16K tokens for MATH500 and AMC23, 32K for AIME24) and freeze it. All probing is then performed *offline* against fixed text prefixes of that frozen trajectory. Because probes never re-enter the decode, changing a probe schedule—or a stopping rule—cannot change the underlying reasoning; every rule is scored against the same trajectories. We build two offline probe banks: a *dense* bank that re-probes every 64 tokens with simple@32 (a 32-sample boxed-answer probe), and an *adaptive* bank that additionally fires event-triggered probes at entropy drops and at conclusion/reflection/answer markers.

3.2 Token accounting

For a rule that stops a problem at token position s after consuming probe output p , we charge total decode tokens $T = s + p$ and compare against the frozen baseline B (the full trajectory, capped at budget). *Net* savings is $(B - T)/B$ and *gross* savings is $(B - s)/B$; their gap is the probe cost. The accuracy drop compares the correctness of the committed answer at s against that of the frozen final answer (it is not an agreement rate between them). Two quantities therefore move independently, and we track both: *where* the rule stops (setting the accuracy drop and the main-token saving s) and *how much it probes* to get there (adding p). Section 5 shows that only the first is intrinsic.

3.3 The consensus rule space (W, s)

The consensus signal in the windowed family we study reduces to two hyperparameters: a **window size** W (how many of the most recent probes to inspect) and a **share threshold** s (what fraction of that window must agree on a single answer before a stop is allowed). This pair spans the family to a close approximation: $W=1$ is the “latest-probe” rule, $s=1.0$ is strict unanimity over the last W probes (the CertIndex-style stop), and an entropy-over-the-window criterion closely tracks the same agreement and adds no materially new operating point on our grid. We sweep $W \in \{1, 3, 5, 8, 12, 16, 24, 30\}$ and $s \in \{0.6, 0.8, 1.0\}$, deliberately including *large* windows so a rule can demand a long, stable agreement before committing.

Around this core we vary only the operational knobs that govern *when* and *how* the signal is read: **probe schedule** (fixed interval $\in \{64, 128, 256, 512\}$ tokens, or event-triggered probing at conclusion/entropy/reflection markers); **maturity** (a minimum-token floor $\in \{0, 512, 1024, 2048, 4096\}$ before any stop is allowed); **validity** (an answer-shape filter, rejecting empty/letter answers or only empties); and **certainty** (optionally requiring the supporting probes to be flagged certain). Enumerating this preregistered grid over a fixed-schedule and an event-triggered family yields 3,520 candidate consensus rules; formal definitions are in Appendix A.

A better signal, swept identically. To test whether the failure is early exit *itself* or the consensus *signal*, we sweep one non-consensus method through the same pipeline, gates, and token accounting. **DEER** (Yang et al., 2025) stops on the model’s own confidence in a trial answer at a reasoning boundary; its only hyperparameter is a confidence threshold, which we sweep over 14 values. We use the stronger *trial-answer-submit* variant (commit the trial answer directly, no separate readout) rather than the faithful readout reproduction—our aim here is not to reproduce prior work but to ask whether *any* rule can clear the gates. Placing DEER and consensus on one frontier under identical accounting is what makes the comparison a controlled test of the signal.

3.4 Models, benchmarks, splits

Development uses two models—DeepSeek-R1-Distill-Qwen-7B and Qwen3-8B—on three bench-

Operating point	Max total acc. drop	Min total saving	psf floor
conservative	1.0 pp	10%	0.80
balanced	2.0 pp	20%	0.80
token_efficient	3.5 pp	30%	0.70

Table 1: Preregistered acceptance gates. A rule must keep the macro-averaged accuracy drop under “max total acc. drop,” deliver at least “min total saving” net tokens, and post positive net saving on at least a psf fraction of environments. Gates are fixed in advance and *not* relaxed post hoc.

marks (MATH500, AMC23, AIME24), each at seeds 42/43/44: 18 model×benchmark×seed environments. Problems are partitioned by *problem id* into 60/20/20 train/dev/test splits (MATH500 300/100/100, AMC23 24/8/8, AIME24 18/6/6), so no problem leaks across splits. Confirmation reserves seeds 45/46/47 plus two unseen models—DeepSeek-R1-Distill-Llama-8B and Distill-Qwen-32B—all evaluated on the *test* split only. In one line: *train* generates and screens candidate rules in-sample, *dev* is the sole selection target the gates are read on, and *test* (with the two unseen models) is read once for confirmation after the rules and thresholds are frozen—never for tuning. Where later sections say “train+dev,” they mean a rule was required to pass a gate in-sample on train and is then measured on dev.

3.5 Preregistered operating points and gates

We fix three operating points and their acceptance gates *before* looking at held-out data (Table 1). Each gate is applied in order: first the **total accuracy drop** (macro-averaged over the 18 environments) must stay under a cap; then the **total net token saving** (macro-averaged) must clear a floor; and finally a **positive-saving fraction** (psf, the fraction of environments with strictly positive net saving) floor must be met. A rule is accepted at an operating point only if it satisfies all three on train+dev. The saving floor is what distinguishes a genuinely useful early exit from a rule that stays accurate only by almost never stopping.

3.6 Metric resolution

The total accuracy drop is a macro-average over the 18 environments (each model×benchmark×seed weighted equally). We span two models, three benchmarks, and three seeds and weight them equally so that selection rewards a rule that is robust across benchmarks, models, and randomness, not

one that overfits a single setting. Pooling problems instead would let MATH500 (100 dev problems per seed vs. 8 for AMC23 and 6 for AIME24) dominate the headline, so a rule that wins only on MATH500 could pass; equal per-environment weight forces a selected rule to hold on the harder competition sets too. Macro also keeps any single problem from swinging the result—one AIME24 problem is 16.7 pp of a 6-problem split but only $\sim 1/108$ of the macro drop—while the three seeds damp the extra per-cell noise of the small sets. The gates are read on the dev split (the selection target) with train reported alongside; because *no* consensus rule clears any gate by a wide margin (Section 4), the finding does not rest on a single fragile cell. Held-out confirmation on the test split and two unseen models provides the complementary out-of-distribution check.

3.7 Preregistration commitments

Three commitments make the result interpretable. (C1) The test split and both confirmation models are never read during rule selection. (C2) All three operating points and their gates in Table 1 are fixed before seeing held-out data and are not relaxed afterwards; if no rule passes a gate, that is reported as the finding rather than engineered away, and all three points are reported regardless of outcome. (C3) The frozen rule set is fixed on train+dev and its hash recorded before any test/confirmation evaluation. Section 4 reports what happened when we applied C2.

4 Experiments

Section 5 argues that probe agreement cannot make a stop safe on its own, because the probes are not independent readings. The experiments here establish that this is not a matter of tuning: we search the rule space exhaustively and under commitment, and find nothing that escapes it.

We first confirm the phenomenon the paper turns on—intermediate consensus is non-terminal (§4.1)—then report the preregistered sweep and its central negative result: *no* consensus rule is both safe and saving (§4.2). We then show the failure is specific to the *signal*, not to early exit, by sweeping a non-consensus method (DEER) through the identical pipeline and gates and finding that it *does* clear them (§4.3); a released consensus stopper, CertIndex (CoT), shows the same failure at its default setting (§4.4). Finally we confirm the frontier

out of distribution (§4.5).

4.1 Non-Terminal Agreement

We log a dense probe stream for every MATH500 problem under DeepSeek-R1-Distill-Qwen-7B at the main 16K/32K budgets (all 500 problems, three seeds; 1,500 trajectories), never intervening—the probe is a separate forward pass that only observes the main trajectory (full setup in Appendix G)—and read off four facts.

(1) Final agreement is reliable; early agreement is not. By the time a trajectory finishes, agreement does track correctness: *cumulative* agreement across all probes is 97.8% correct (186 trajectories reach it), and even the final five-probe *window* is 90.4% correct (1,205). But an online rule cannot wait for the end—it fires the *first* time probes agree, mid-trajectory, where the answer is far from settled (facts 2–4). The operative question is not whether *final* consensus is trustworthy, but whether *early* consensus is—and it is not. **(2) Recovery is common.** Of the 1,148 trajectories whose *first* probe is wrong, 89.1% eventually flip to correct; of the 736 that formed a three-probe consensus *different* from their final answer, 84.2% ended up correct. **(3) Later consensus is worse.** Consensus formed before 512 tokens is 91.0% correct, but consensus forming only after 2048 tokens is 71.6% correct—hard problems converge slowly *and* to worse answers. **(4) The cost lands at the stop.** A naive stop (first three consecutive agreeing, certain, non-empty probes) fires on 1,477/1,500 trajectories and commits to answers correct only 50.5% of the time, whereas letting those same problems run to completion reaches 90.7%: a 40.2 pp loss.²

The cost is concentrated early (Figure 4, appendix). Expressing the moment consensus first forms as a *fraction* of each trajectory’s own length—which controls for problem difficulty—half of all consensuses form within the first tenth of the trajectory, and there the answer they commit to is correct only 27.5% of the time against 85.2% for the same problems run to completion. The gap narrows as consensus forms later, but an online rule cannot choose to wait: it fires when agreement first appears, which is precisely where agreement is least informative.

²A simplified heuristic, *not* the CertIndex (CoT) rule (§4.4); it fires early because the first stable, certain agreement is typically a mid-trajectory placeholder. The swept rules below and the mechanism of Section 5 quantify the same directional accuracy tax across the full rule space.

A good stopping rule must therefore be much more than “agreement.” The rest of this section asks whether the *best possible* such rule, drawn from a large preregistered space, can be both safe and economical.

4.2 Sweep Results

We replay all 3,520 consensus rules on the frozen trajectories and aggregate to per-environment metrics—126,720 rows = 3,520 rules $\times$ 18 environments $\times$ 2 splits (train, dev)—then, per rule, its total (macro) accuracy drop and net token saving (accounting of Section 3.2). Applying the preregistered gates yields the negative result that Section 5 explains.

The outcome holds across environments: across all 126,720 rule–environment evaluations accuracy strictly *falls* in 64.9% and strictly *rises* in only 10.5% (unchanged in 24.6%), for a mean drop of 13.0 pp; over rules the total dev drop has median 13.2 pp and a minimum of 0.81 pp. **Not one of the 3,520 rules clears any gate.** The reason is the shape of the trade-off, read directly off the sweep’s own two axes (Figure 5, appendix): staying accurate and saving tokens are in direct tension. Only 40 rules keep the total drop at or below the conservative 1.0 pp cap, and the most any of them saves is 0.2%—far under the 10% floor; conversely, the first rule to save 10% already costs 2.66 pp, saving 20% costs 6.17 pp, and saving 30% costs 11.8 pp. Crucially, the sweep now includes *large* windows (up to $W=30$): these do buy low accuracy drop, but only by stopping so late that net saving collapses to zero. No setting of the consensus signal escapes the trade-off: read cell by cell across the whole (W, s) surface, the safest rule available at each cell approaches the 1 pp cap only as the saving it buys approaches zero, while every cell that saves substantially costs several points, and no cell clears the conservative gate (Figure 5, appendix).

4.3 A Non-Consensus Signal

The same gates are cleared, comfortably, by a rule that does *not* read probe agreement. Sweeping DEER’s confidence threshold through the identical pipeline and token accounting (Section 3.3), its trial-answer-submit variant passes all three operating points on dev: the conservative point loses only 0.33 pp while saving 28.2%, the balanced point 1.03 pp at 29.6%, and the token-efficient point 2.75 pp at 31.9%; at a stricter threshold it is accuracy-neutral (−0.06 pp) while still saving

Method	Acc.	Δ acc (macro)	Net saving
Full generation	82.5%	—	0%
Consensus (drop $\leq$ 1.0)	81.6%	−0.9 pp	+0.2%
Consensus (save $\geq$ 10%)	79.9%	−2.7 pp	+10.7%
DEER (conservative)	82.5%	−0.33 pp	+28.2%
DEER (balanced)	81.9%	−1.03 pp	+29.6%
DEER (token_eff.)	80.1%	−2.75 pp	+31.9%

Table 2: Dev operating points (macro over 18 environments). No consensus rule is simultaneously safe (≤ 1.0 pp) and saving ($\geq 10\%$): capping the drop forces saving to near zero, and demanding saving forces a large drop. DEER, reading a non-consensus signal, clears all three gates on the same trajectories under the same accounting.

20.8%. Early exit is therefore not impossible; the consensus *signal* is what cannot make it safe (Table 2).

4.4 CertaIndex (CoT) at Its Default

The sweep covers rules we constructed; a released method lets us see where a default setting lands. We reproduce two prior stoppers in a common harness—the same frozen trajectories, the same token accounting. **CertaIndex (CoT)** (Fu et al., 2024) is the consensus stop for a single chain of thought, reproduced from the released implementation at its default mid setting: stop once three consecutive probes give the same non-empty answer and none hedges. We use the authors’ probe prompt and probe schedule, so this is their rule as shipped, not a variant of ours. **DEER** (Yang et al., 2025) here is the readout reproduction (the stronger trial-answer-submit variant is what we swept in §4.3).

Two scope notes. First, CertaIndex’s multi-path setting is ordinary self-consistency: it measures entropy over answers from *independently sampled* trajectories, and nothing here bears on it. What we reproduce is the single-trajectory case, where the same statistic is read over repeated probes of one chain. Second, Table 3 reports these reproductions on our frozen trajectories and benchmarks; it is not a re-run of either paper’s end-to-end system, and the original CoT results were reported on an easier benchmark, where intermediate answers settle sooner.

At a 99.8% stop rate and 70 pp lost, CertaIndex (CoT) lands where the swept consensus frontier says a rule at that stop rate must land. This is a property of the signal, not a defect of the imple-

mentation: it is the same trade-off §4.2 maps across the whole rule space, seen at one shipped operating point. The DEER readout, reading a non-consensus signal on the same trajectories, stays near full accuracy; the stronger trial-answer-submit variant we swept turns that into a gate-clearing frontier (§4.3). That only the non-consensus signal survives is what we analyze in Section 5.

4.5 Generalization

We read the test split once (commitment C3). Because the development result is that *no* consensus rule clears any gate, there is nothing to confirm by re-selection; we ask whether the *frontier* reproduces, ranking rules on dev and measuring on test.

Held-out test split. Each consensus rule’s total drop on test tracks its dev value closely ($r = 0.98$). The split-invariant claim is the empty *joint* gate: **no consensus rule clears the conservative gate on both splits**—0 on dev, and the 444 rules admissible on *test alone* are exactly the in-sample winners the held-out split exists to reject (all fail on dev). DEER, by contrast, remains admissible across splits: 3 of its thresholds clear the conservative gate on *both* dev and test (and 5 clear balanced on both), so the cross-signal contrast is not a dev artifact.

Unseen scale (4 $\times$) and architecture. Both unseen models are evaluated on the test split at three seeds (45/46/47; nine environments each), so their gates are as well resolved as development’s. On Distill-Qwen-32B, never used in development, the consensus frontier reproduces ($r = 0.97$ against dev drop); on the held-out architecture and family DeepSeek-R1-Distill-Llama-8B it reproduces at $r = 0.94$.³ The *conservative* gate stays empty on every model—dev, 32B, and Llama all admit 0 rules (the best rule under a 1.0 pp drop saves 0.6% on 32B and 9.3% on Llama). At the looser gates a scale effect appears: on the 4 $\times$ larger 32B a few consensus rules become admissible in-sample (4 at balanced, 6 at token-efficient), reflecting that a bigger model’s intermediate answers stabilize somewhat earlier—but these are not the dev-selected rules (dev admits none), and on the same-scale Llama architecture the gates stay empty (0/0/0). Consensus thus yields no rule that is safe-and-

³The Llama-8B trajectories were re-collected after correcting a chat-template defect (a missing model-specific beginning-of-sequence token that degraded its generations); the corrected model scores $\sim 88\%$ on MATH500.

Method	Model	Accuracy	Δacc	Fair token saving	Stop rate	Signal
Full generation	Qwen3-8B	85.4%	—	0%	0%	—
Full generation	DeepSeek-R1-Distill-Qwen-7B	79.8%	—	0%	0%	—
CertaIndex (CoT)	Qwen3-8B	15.3%	−70.1 pp	90.1%	99.8%	consensus
CertaIndex (CoT)	DeepSeek-R1-Distill-Qwen-7B	23.9%	−55.9 pp	76.7%	98.7%	consensus
DEER	Qwen3-8B	86.2%	+0.8 pp	16.3%	41.4%	boundary conf.
DEER	DeepSeek-R1-Distill-Qwen-7B	74.9%	−4.8 pp	20.2%	56.1%	boundary conf.

Table 3: Prior early-stoppers reproduced on our frozen trajectories (dev, macro over three benchmarks; probe tokens charged). **CertaIndex (CoT)** stops on almost every problem and stops early, and loses 56–70 pp while saving 77–90% of tokens—the pattern §4.1 describes, at a single shipped operating point. **DEER**, reading boundary confidence rather than probe agreement, stays near full accuracy (+0.8 pp on Qwen3-8B) while still saving tokens. These are reproductions on our benchmarks, not the original papers’ end-to-end numbers; Section 5 interprets the contrast.

saving where it is selected, on any model. DEER, swept identically on these unseen models, **clears the gates on both**: on Qwen-32B it is accuracy-neutral at the conservative point (−0.24 pp while saving 32.4%, $\tau=0.97$) and on Llama-8B it loses 0.67 pp at 26.7% ($\tau=0.99$), passing all three gates. Consensus is unsafe, and DEER safe-and-saving, across all four axes we vary—benchmark, seed, 4× scale, and architecture/family.

The selection pipeline runs end to end in Figure 6 (appendix). The 478 consensus rules that clear the conservative gate in-sample on train are tracked onto dev and test as a band; the three dev-selected DEER thresholds are tracked individually. On dev—the split the gate is applied on—not one of the 478 is admissible, and the joint gate is empty, while all three DEER points stay inside on every split. The dev and test splits are not equally demanding: the same 478 rules have a median drop of 4.5 pp on dev but 0.6 pp on test, an asymmetry driven by the two smallest environments (AMC23 and AIME24 under Qwen3-8B, 8 and 6 problems), which carry the same weight as a 100-problem environment under macro-averaging. This is why the claim is stated jointly rather than per split: 364 of the 478 are admissible on test, but they are exactly the in-sample winners a held-out split exists to reject, and none of them is selectable from dev.

Per-model and per-benchmark views of the same data, together with an *oracle* upper bound—a perfect probe-based stop committing the earliest correct probe, which sits at a negative drop of 2–5 pp at 40–80% saving, far in a corner neither method reaches—are given in Appendix E.

5 Analysis

Section 4 shows that no rule in the searched space is both safe and saving. This section gives the reason, which is what makes the result more than a report on 3,520 particular rules: a *stability–terminality gap*. Under a fixed probing procedure, agreement establishes that the current answer *persists*, not that the reasoning has *terminated*. The source is that a probe does not read a settled belief—when the reasoning is unfinished it *forces* an answer, and the model’s cheapest output is a repeat of its last one, an elicited placeholder that can look stable while the trajectory keeps correcting itself. We establish this directly, and in this order: the answer a probe returns is sensitive to how the probe is worded early on (§5.1); a hand-labelled taxonomy shows most stopped-but-wrong commitments are to answers the model had not converged on (§5.2); and the correction that follows makes the cost *directional*, so widening the window—the obvious remedy—reduces the loss only by declining to stop (§5.3). A structural precondition underlies all three: probe readings are not independent samples, so agreement supplies no vote that could catch the error (§5.7). The resulting recovery asymmetry is a property of the reasoning trajectory, not of a particular threshold.

5.1 Probe Wording versus Position

If early agreement reflected a settled belief, the answer a probe returns should not depend much on how the probe is worded. We test this directly on paired re-probes of the *same frozen prefixes*: at every probe position we read the trajectory with two different probe wordings (our boxed-answer suffix and the CertaIndex suffix, both capped at 32 tokens), so the model’s state is identical and only

the elicitation differs.

Figure 2(a) shows the result. In the first tenth of a trajectory the two wordings return different answers 57% of the time; by the final third they disagree only 11% of the time. Early “answers” are therefore substantially a property of the question asked rather than of the state being read, and they become properties of the state only as the trajectory finishes.

This distinguishes our account from the obvious alternative. If the model simply held a stable but mistaken belief early on, both wordings would read back the *same* wrong value; instead they disagree. What a consensus rule treats as a converging answer is, early on, an output the probe has elicited. Note also that probe wording has almost no effect on *final* accuracy in the same paired data (a 0.65 pp readout effect against a 9.15 pp effect of *when* one probes), so the sensitivity above is specific to early positions and is not a defect of any particular suffix.

5.2 Error Taxonomy

Two annotators independently labelled all 134 stopped-but-wrong cases (Figure 2(b), $\kappa = 0.82$): only one in five is a case where the model had genuinely settled on a wrong value. Three in five are answers it had not converged on at all, and a further one in six are artifacts of the probe’s output format.

The reason is that a probe does not read the model’s belief; it *forces* an answer. A model that has not finished reasoning must still write something in the box, and the cheapest thing to write is what it wrote last time. Repetition of a forced answer is therefore stable for reasons that have nothing to do with the reasoning being finished, and a consensus rule cannot tell that case apart from a settled one. This also predicts the window effect of §5.3: a longer window is a longer buffer in which a placeholder can be overturned, which is why it reduces false consensus without removing the underlying risk.

5.3 Effect of Window Size

The natural remedy is to demand more agreement before stopping. It helps, but only by stopping less often. Consider only the stops where committing early *changes* correctness relative to the trajectory’s own final answer, and count two outcomes: (final-correct, stop-wrong)—a correct recovery destroyed—versus (final-wrong, stop-correct)—a wrong answer luckily banked. Pure sampling noise predicts a $\approx 1:1$ ratio; a genuine “continued rea-

soning corrects” effect predicts $\gg 1$. On the frozen replay this ratio depends strongly on the window: it falls from $\approx 45:1$ for a latest-probe stop ($W=1$) to $\approx 2:1$ at the largest window ($W=30$). But the low end is bought only by stopping so rarely that net saving vanishes ($W=30$ fires on 121 dev problems vs. 668 at $W=1$); at any window that actually saves tokens the ratio stays well above 1:1. DEER, by contrast, holds $\approx 2.4\text{--}3.5:1$ at its three operating points *while* saving 28–32%, because its boundary-confidence signal commits mainly once the trajectory has settled. Stopping on consensus therefore removes a correct-in-the-end answer many times more often than it rescues a wrong one—directional, not sampling noise. This is the recovery phenomenon of §4.1 (84.2%, 89.1%), now at the main 16K/32K budgets and as a signed effect. Figure 3 plots the whole window axis, so the claim rests on the shape of the curve rather than on either endpoint: the ratio and the saving fall together, and DEER sits at consensus’s best ratio and its best saving simultaneously, which no window does.

5.4 Case Studies

The three regimes above appear together on individual trajectories (DeepSeek-R1-Distill-Qwen-7B, MATH500, probes every 64 tokens; $a \times n$ denotes an answer repeated over n consecutive probes). Appendix F gives the full streams and discussion.

On problem 320 the model answers $\emptyset$ at its *first* probe and holds it for 27 probes—nearly half a 3,683-token trajectory—before reasoning its way to the correct $-\frac{24}{25}$. The committed value is not a considered wrong answer but one emitted before any real work, and its stability is correspondingly uninformative. The run is also long enough to defeat every window we searched: even $W=24$ fires here and commits $\emptyset$, which is why widening the window is not a remedy for a placeholder that outlasts it.

Problem 253 asks what fraction of two feet is three inches, a free-response question with no options. Every probe from token 384 to token 1600 returns the letter D; the correct $\frac{1}{8}$ appears only at the final probe. Agreement is total and refers to nothing, which is the format artifact of §5.2 in its clearest form.

The minority case is worth stating too. On problem 240 a three-probe rule commits 52 while the trajectory itself ends at 154; the gold answer is 116, so both are wrong and stopping destroys nothing. This is the regime consensus early exit was




Figure 2: **Two sides of the stability-terminality gap.** (a) The same frozen prefix (DeepSeek-R1-Distill-Qwen-7B, MATH500) re-probed with two differently worded answer queries (our boxed-answer suffix and the CertIndex suffix, both capped at 32 tokens), so the model’s state is identical and only the elicitation differs. Early in a trajectory the two wordings return the *same* current answer less than half the time (43% in the first tenth, so they *disagree* 57% of the time), and the probed answer is usually wrong; by the end they nearly always coincide (89%) and are usually correct. Position is a fraction of each trajectory’s own length. (b) All 134 stopped-but-wrong cases, labelled independently by two annotators ($\kappa=0.82$ on the coarse coding). Three in five commit to an answer the model had *not converged* on, and a further one in six to a probe-format artifact; fewer than one in five are cases where it had genuinely settled on a wrong value.

designed for—the model was not going to solve the problem, so a stop is free. The difficulty is that agreement cannot tell this case apart from problem 320, and problem 320 is the more common one.

5.5 Why the recovery asymmetry is trajectory-level

Why does the ratio stay above one across the grid? The recovery it counts is a property of the reasoning trajectory: for a stop at token position s , the correction forgone is whatever the trajectory would have made after s , which depends on the trajectory, not on how it was probed. The *realized* accuracy cost of a rule is not itself probe-independent—it depends on where the rule stops, and hence on the probe schedule and stopping rule—but the recovery asymmetry that produces it is present at every position where reasoning is still moving. Our probe grid (intervals 64/128/256/512 tokens) makes essentially every stop position reachable, so the floor we observe is the minimum over these schedules rather than a lucky miss; a sparser schedule can only stop later or at a subset of positions, never at a safer one the grid missed. Hence no threshold on the consensus signal *within the searched space* escapes the tax. This does not turn a finite rule search into a universal impossibility theorem—it motivates caution for consensus variants beyond the preregistered schema rather than proving a uni-

Consensus point	Gross saving ($B - s$)/ B	Net saving ($B - s - p$)/ B
save $\geq 10\%$ (drop 2.7 pp)	+14.7%	+10.7%
save $\geq 20\%$ (drop 6.2 pp)	+27.0%	+20.2%

Table 4: Gross vs. net savings for two consensus frontier points; the gap ($\sim 4-7$ pp) is the probe tax. It shifts the *savings* axis but never rescues the accuracy drop—which is the intrinsic tax that keeps the corner empty.

versal claim about them. In short, agreement reports what the model has said recently, not whether its reasoning has finished.

5.6 Accuracy Tax versus Probe Tax

The natural rebuttal is that the negative result is an artifact of charging dense re-probing. It is not, and separating two costs shows why. Besides the accuracy tax, a rule pays a *probe tax*: the probe output p (Section 3.2). Under dense re-probing a late-stopping large-window rule pays p over a long prefix, opening a gap between its *gross* savings and its (smaller) *net* savings (Table 4). Unlike the accuracy tax, the probe tax scales with probe density and is reducible (sparser or shorter probes, KV-cache reuse). The two are separable: the *savings* axis is partly a probe-tax artifact and is addressable, but the *accuracy* floor is intrinsic to the trajectory—present wherever reasoning is still moving—and is what makes the corner empty; it is that floor, not the probe tax, that DEER escapes.




Figure 3: **Lowering the ratio means refusing to fire.** Among stops whose committed answer differs in correctness from the trajectory’s frozen final answer, we count recoveries destroyed against wrong answers banked; sampling noise predicts 1:1. Widening the window brings the ratio down from 45:1 to 2:1, but only by stopping on ever fewer problems (668 $\rightarrow$ 121 of 684), so net saving collapses from 92% to 8% alongside it. Dev split, 18 environments; one consensus rule per W (share threshold 1.0, probe interval 128, maturity floor 512, schema validity), with W the only axis that varies.

5.7 The role of independence

A structural fact underlies why a bad stop is never caught. Self-consistency’s majority vote can absorb an error because its paths are sampled independently (§2): a wrong answer on one path is outvoted by others that reasoned differently. Probing one trajectory yields no such thing—every probe reads the same prefix, so the readings move together, and when the prefix is wrong they are wrong together. Agreement across probes is thus a statement about the stability of one line of reasoning, not corroboration from several, so a consensus rule holds no vote in reserve to overturn the placeholder it commits to. Independence explains why the gap is *dangerous* rather than merely present; it does not on its own show the committed answer is non-terminal—that is what the direct measurements above establish.

5.8 Locating the Failure

The two experiments of §4.3–4.4 localize the cause. Every method that stops on *probe agreement*—the swept consensus family and the CertIndex (CoT) reproduction alike—fails a gate or loses many points; the one method that clears the gates, DEER, reads not probe agreement but the model’s own confidence in a trial answer at a reasoning

boundary. Because DEER is swept through the *identical* pipeline, gates, and token accounting, this is a controlled test: the same early-exit machinery fails on the consensus signal and succeeds on a non-consensus one, so the failure is a property of the *signal*, not of early exit.

What this does *not* establish is which signal should replace agreement. We show that agreement alone is insufficient; we do not show that any particular alternative is sufficient, and DEER is a control rather than a recommendation. Identifying and validating a safe stopping signal—with its own train/dev $\rightarrow$ test confirmation—is left open.

That qualification matters. Our claim is about consensus used as the sole criterion for stopping, which is the use we searched and the use that fails. Probe agreement is cheap to compute and does carry one piece of information—that the current answer has not changed—and combined with a signal that reflects the model’s own estimate it may well be useful; CertIndex composes multiple indicators for related reasons (Fu et al., 2024). Nothing here argues against such combinations. What the sweep rules out is the step of treating agreement, on its own, as sufficient evidence that reasoning has finished.

Two further caveats bound the claim: it concerns *consensus-based* rules within the searched schema and dense-probe regime, and it does not say early exit is impossible—only that this signal cannot make it safe alone.

6 Conclusion

We asked whether intermediate answer-consensus is a safe signal on which to early-exit a reasoning LLM. Used on its own it is not, and not because the rule is untuned: under a fixed probing procedure agreement establishes that the current answer *persists*, not that the reasoning has *terminated*—a *stability–terminality gap*. What the probes agree on is often an answer the model has not settled on—three in five of our hand-labelled wrong stops commit to a placeholder the probe forced it to produce, and re-wording the probe changes the early answer more often than not—because a probe forces an output rather than reading a finished belief. Stopping on such an answer destroys a correct-in-the-end answer far more often than it rescues a wrong one (up to $\sim 45\times$ at an aggressive latest-probe stop, $\sim 2\times$ at the largest windows, where the rule has nearly stopped firing), so demanding more agreement low-

ers the loss only by declining to fire. The recovery asymmetry behind this is a property of the trajectory, and it shows up as an empty safe-and-saving corner across a preregistered space of 3,520 rules—not one of which clears any gate—reproduced on a held-out split and two unseen models, and visible at one shipped operating point in CertIndex (CoT), which loses 56–70 pp. That the corner is empty for consensus but not for early exit is a positive control, not a proposal: swept through the identical protocol, a method reading the model’s own boundary confidence (DEER) clears all three gates on the very same trajectories, so the failure is specific to consensus rather than to early exit. This does not put probe agreement beyond use—it is cheap, and it does report that the current answer has not changed; what our results rule out is treating that report, by itself, as evidence that the reasoning is done. Repeated agreement under a fixed probing procedure establishes answer persistence, not reasoning terminality; within the broad preregistered family studied here, making that agreement stricter only trades away stopping opportunities rather than turning it into a safe stopping signal. Agreement fails not because it is insufficiently strict, but because it repeatedly measures the wrong object.

Limitations

The mechanism is argued, not isolated experimentally. Our account of *why* consensus fails—that probes of one trajectory are not independent readings, and that they frequently record a forced placeholder—rests on a structural argument (Section 2) together with the hand-labelled error taxonomy and the window sweep. We do not run a controlled experiment that manipulates independence directly, and the labelled cases were collected under a short window, the regime where unconverged answers are most common. A stronger test would construct probe streams with varying degrees of dependence and measure how the accuracy tax responds.

Scope of the negative result. The central result—no consensus rule clears any gate—is established on 18 development environments (two models, three benchmarks, seeds 42/43/44) and confirmed on the held-out test split and two unseen models (Section 4). It is a statement about the space searched: one boxed-answer probe suffix (simple@32) and one preregistered consensus schema (3,520 rules over a window-size $\times$ share-threshold

family plus operational knobs). Other probe wordings and signals outside this schema are not searched; the mechanism (Section 5) is our argument that consensus variants would not help, but it is an argument, not an exhaustive proof.

DEER as the positive control, not a benchmarked SOTA. We sweep DEER’s trial-answer-submit variant to establish that *some* signal clears the gates the consensus family fails (on the development models, the held-out test split, and both unseen models); we do not claim it is the best possible early-exit method. Its role is a controlled positive control for the signal, not a leaderboard entry.

Small held-out sets on hard benchmarks. AIME24 and AMC23 are small (dev splits of 6 and 8 problems per seed), and the unseen models contribute a single seed, so those per-environment cells are noisy. The total drop macro-averages over the 18 development environments (three seeds per benchmark $\times$ model), and we read the single-seed unseen models as frontier-reproduction evidence rather than as independent gate tests (Section 3.6).

Probe-tax dependence of the savings axis. Net-savings numbers charge dense re-probing (every 64 tokens, 32-token probes); a sparser or KV-cache-reusing schedule would change the *savings* axis and could move some rules to positive net savings. We separate gross from net savings (Section 5) precisely so the probe-tax-dependent part is distinguishable; the accuracy-tax result does not depend on the schedule.

Domain. All benchmarks are competition mathematics with checkable final answers. Whether false consensus and the accuracy tax transfer to open-ended reasoning, code, or agentic tasks is out of scope.

References

- Pranjal Aggarwal, Aman Madaan, Yiming Yang, and Mausam. 2023. Let’s sample step by step: Adaptive-consistency for efficient reasoning and coding with LLMs. In *Proceedings of EMNLP*.
- Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré, and Azalia Mirhoseini. 2024. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv preprint arXiv:2407.21787*.

987	Xingyu Chen, Jiahao Xu, Tian Liang, Zhiwei He,	Qwen Team. 2025. Qwen3 technical report . <i>arXiv preprint arXiv:2505.09388</i> .	1042
988	Jianhui Pang, Dian Yu, Linfeng Song, Qiuzhi		1043
989	Liu, Mengfei Zhou, Zhuosheng Zhang, Rui Wang,		
990	Zhaopeng Tu, Haitao Mi, and Dong Yu. 2024. Do	Lee Ross, David Greene, and Pamela House. 1977. The	1044
991	not think that much for $2+3=?$ on the overthinking of	“false consensus effect”: An egocentric bias in so-	1045
992	o1-like LLMs. <i>arXiv preprint arXiv:2412.21187</i> .	cial perception and attribution processes. <i>Journal of</i>	1046
		<i>Experimental Social Psychology</i> , 13(3):279–301.	1047
993	DeepSeek-AI. 2025. DeepSeek-R1: Incentivizing rea-	Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani,	1048
994	soning capability in LLMs via reinforcement learning .	Dara Bahri, Vinh Q. Tran, Yi Tay, and Donald Met-	1049
995	<i>arXiv preprint arXiv:2501.12948</i> .	zler. 2022. Confident adaptive language modeling.	1050
996	Zhe Dong, Fang Qin, and Manish Shah. 2026. When	In <i>Advances in Neural Information Processing Sys-</i>	1051
997	does learning to stop help? a cost-aware study of	<i>tems (NeurIPS)</i> .	1052
998	early exits in reasoning models. <i>arXiv preprint</i>		
999	<i>arXiv:2606.30852</i> .	Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Ku-	1053
1000	Yichao Fu, Junda Chen, Siqi Zhu, Zheyu Fu, Zhong-	mar. 2024. Scaling LLM test-time compute optimally	1054
1001	dongming Dai, Aurick Qiao, and Hao Zhang. 2024.	can be more effective than scaling model parameters.	1055
1002	Efficiently serving LLM reasoning programs with	<i>arXiv preprint arXiv:2408.03314</i> .	1056
1003	Certainindex . <i>arXiv preprint arXiv:2412.20993</i> .	Peiyi Wang, Lei Li, Zhihong Shao, Runxin Xu, Damai	1057
1004	Alex Graves. 2016. Adaptive computation time	Dai, Yifei Li, Deli Chen, Yu Wu, and Zhifang Sui.	1058
1005	for recurrent neural networks. <i>arXiv preprint</i>	2024. Math-shepherd: Verify and reinforce LLMs	1059
1006	<i>arXiv:1603.08983</i> .	step-by-step without human annotations. In <i>Proce-</i>	1060
1007	Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul	<i>ings of ACL</i> .	1061
1008	Arora, Steven Basart, Eric Tang, Dawn Song, and	Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le,	1062
1009	Jacob Steinhardt. 2021. Measuring mathematical	Ed Chi, Sharan Narang, Aakanksha Chowdhery, and	1063
1010	problem solving with the MATH dataset. In <i>NeurIPS</i>	Denny Zhou. 2023. Self-consistency improves chain	1064
1011	<i>Datasets and Benchmarks Track</i> .	of thought reasoning in language models. In <i>Inter-</i>	1065
1012	Parsa Hosseini, Sumit Nawathe, Mahdi Salmani,	<i>national Conference on Learning Representations</i>	1066
1013	Meisam Razaviyayn, and Soheil Feizi. 2026. Early	(ICLR).	1067
1014	stopping for large reasoning models via confidence	Chenxu Yang, Qingyi Si, Yongjie Duan, Zheliang Zhu,	1068
1015	dynamics. <i>arXiv preprint arXiv:2604.04930</i> .	Chenyu Zhu, Zheng Lin, Li Cao, and Weiping Wang.	1069
1016	Jie Huang, Xinyun Chen, Swaroop Mishra,	2025. Dynamic early exit in reasoning models . <i>arXiv</i>	1070
1017	Huaixiu Steven Zheng, Adams Wei Yu, Xiny-	<i>preprint arXiv:2504.15895</i> .	1071
1018	ing Song, and Denny Zhou. 2024. Large language	A Rule Schema Details	1072
1019	models cannot self-correct reasoning yet. In <i>Inter-</i>	The consensus schema of Section 3 centers on	1073
1020	<i>national Conference on Learning Representations</i>	two hyperparameters and a few operational knobs,	1074
1021	(ICLR).	enumerating to 3,520 candidate rules over a fixed-	1075
1022	Yiwei Li, Peiwen Yuan, Shaoxiong Feng, Boyuan	schedule and an event-triggered family. Each rule	1076
1023	Pan, Xinglin Wang, Bin Sun, Heda Wang, and Kan	fixes:	1077
1024	Li. 2024. Escape sky-high cost: Early-stopping	1. Window size $W \in \{1, 3, 5, 8, 12, 16, 24, 30\}$:	1078
1025	self-consistency for multi-step reasoning. In <i>Inter-</i>	the number of most recent probes inspected.	1079
1026	<i>national Conference on Learning Representations</i>	$W=1$ is the latest-probe rule.	1080
1027	(ICLR).	2. Share threshold $s \in \{0.6, 0.8, 1.0\}$: the frac-	1081
1028	Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri	tion of the last W probes that must agree on	1082
1029	Edwards, Bowen Baker, Teddy Lee, Jan Leike, John	one answer (share measured over the full win-	1083
1030	Schulman, Ilya Sutskever, and Karl Cobbe. 2024.	indow, so an empty or dissenting probe counts	1084
1031	Let’s verify step by step. In <i>International Conference</i>	against consensus); $s=1.0$ is strict unanimity.	1085
1032	<i>on Learning Representations (ICLR)</i> .	An entropy-over-the-window criterion closely	1086
1033	Xin Liu and Lu Wang. 2025. Answer convergence as	tracks the same agreement and adds no materi-	1087
1034	a signal for early stopping in reasoning . In <i>Proce-</i>	ally new operating point on our grid, so it is	1088
1035	<i>edings of the 2025 Conference on Empirical Methods</i>	not swept separately.	1089
1036	<i>in Natural Language Processing (EMNLP)</i> .	3. Probe schedule: fixed interval	1090
1037	Brian A. Nosek et al. 2025. Preregistration in machine	$\in \{64, 128, 256, 512\}$ tokens, or	1091
1038	learning research. <i>arXiv preprint</i> .		
1039	OpenAI. 2024. Learning to reason with		
1040	LLMs. https://openai.com/index/		
1041	learning-to-reason-with-llms/ .		

adaptive_event with entropy-drop and marker triggers and a fallback interval.

4. **Maturity:** a minimum-token floor $\in \{0, 512, 1024, 2048, 4096\}$ before any stop is allowed.
5. **Certainty:** optionally require the supporting probes flagged certain (minimum fraction).
6. **Validity:** answer-shape filter (reject empty / single-letter answers to suppress the Type-E probe artifact of Appendix G), or accept any non-empty answer.

Certainty. A probe is *certain* when its boxed answer is emitted with high token-level confidence: concretely, the mean probability of the answer tokens under the probe forward pass exceeds a fixed threshold (the Certaindex certainty flag). For the simple@32 bank this is aggregated over the 32 samples as the fraction agreeing on the modal answer. Dimension 5 gates a stop on the fraction of the supporting probes that are certain.

Entropy triggers. The adaptive_event schedule fires an extra probe when the token-level predictive entropy of the main decode drops by at least ΔH relative to a trailing baseline (an “the model just committed” signal), and at conclusion/reflection/answer marker tokens (e.g., “therefore”, “the answer is”), subject to a minimum inter-probe gap and a fallback interval.

B Preregistration Text

The protocol version, split manifest hash, candidate-rule generation script, and the three operating-point gates (Table 1) were fixed before any held-out evaluation. Commitments C1–C3 (Section 3) were recorded with the protocol. All three operating-point gates were fixed in advance; when the conservative gate proved empty we followed C2 and report the empty gate as the finding rather than relaxing it, and all three points are reported for both consensus and DEER regardless of outcome.

C Reproducibility

Frozen trajectories, offline probe banks, the sweep archive (126,720 dev metric rows plus the DEER threshold sweep), and the aggregation scripts are released. Table 5 and Figure 5 were reproduced to the reported precision directly from the released sweep

Threshold τ	Total drop (pp)	Net saving
0.9999	−0.06	20.8%
0.999	0.06	25.4%
0.995	0.33	28.2%
0.99	1.03	29.6%
0.97	2.75	31.9%
0.95	5.87	34.8%
0.90	11.43	42.6%

Table 5: DEER trial-answer-submit threshold frontier on dev (macro over 18 environments, same token accounting). Unlike consensus, low drop coincides with substantial saving—the conservative/balanced/token-efficient gates are all cleared around $\tau \in [0.97, 0.995]$.

archive. Grading uses a robust answer-equivalence check that tries the raw and normalized reference forms and both string and math-expression equality; per-problem correctness for the frozen baseline is recomputed at replay time rather than trusted from collection.

D DEER Threshold Frontier

Figure 5 traces the consensus trade-off along both sweep axes; read down the window axis, the minimum attainable accuracy drop falls monotonically (from 22 pp at $W=1$ to 0.8 pp at $W=30$) while the maximum net saving falls with it (from 98% to 39%), so the low-drop region is reachable only where saving has already collapsed. Table 5 gives the DEER threshold frontier that does reach the gate region.

E Supporting Figures: Sweep Surface, Selection, and Frontiers

This appendix collects the supporting figures referenced from the main text. Figure 4 shows where consensus first forms as a fraction of each trajectory’s length (§4.1); Figure 5 reads the empty safe-and-saving corner directly off the sweep’s (W, s) surface (§4.2); and Figure 6 traces the train $\rightarrow$ dev $\rightarrow$ test selection pipeline (§4.5). Figures 7–9 then give the saving-versus-drop view of the same selection and generalization results, resolved per split, per model and per benchmark, with the oracle upper bound on each panel. They add no claim beyond the main text; they show the distributions the summaries are drawn from.

F Case Studies

The mechanism of Section 5 is easiest to see on individual trajectories. All three below are DeepSeek-R1-Distill-Qwen-7B on MATH500, seed 42,




Figure 4: **Consensus forms early, where it is least reliable.** For each dev-split trajectory we take the first window of three agreeing non-empty probes and ask two things: is that answer correct, and would the same problem have been correct if left to run? Position is a fraction of the trajectory’s own length, so 500 tokens into a 600-token trajectory counts as late and 500 into a 16K one does not. Consensus forms in 679 of 684 trajectories, half within the first tenth; the shaded gap is the accuracy a stop at that moment throws away. Open markers mark bins with fewer than 20 trajectories; the final bin ($n=3$) should not be read.

probed every 64 generated tokens; we write $a \times n$ for an answer repeated over n consecutive probes. A rule requiring three agreeing probes fires at the third element of the first such run.

Case 1: a placeholder that outlasts every window we searched. Problem 320 asks for $\cos \angle RPS$ given $\sin \angle RPQ = \frac{7}{25}$ (answer $-\frac{24}{25}$). The probe stream is

$$\begin{aligned} &0 \times 27, B \times 8, -24/25, B \times 7, \\ &E, B \times 6, -24/25 \times 7 \end{aligned}$$

The model answers 0 at the very first probe and repeats it for 27 consecutive probes, covering tokens 64–1728 of a 3,683-token trajectory. It then passes through a stretch of bare option letters before reaching $-\frac{24}{25}$ and holding it to the end; the trajectory is correct.

This single case makes three points. The stop is not a considered wrong answer but a value emitted before any real work has been done, and it is perfectly stable for nearly half the trajectory. Its stability therefore says nothing about whether reasoning has finished. And because the run is 27 probes long, it defeats *every* window in our sweep: even $W=24$, the second-largest we searched, fires here and commits 0. Enlarging the window does

not address a placeholder that persists longer than the window.

Case 2: the probe’s format, not the model’s answer. Problem 253 is “What fraction of 2 feet is 3 inches?”—a free-response question with no answer options. The stream is

$$3, B, 3, 24, B, D \times 20, 1/8$$

From token 384 to token 1600, more than two thirds of the trajectory, every probe returns the letter D. There is nothing for D to refer to. The model is not reporting a belief about the answer; it is completing the probe suffix in a plausible format while its reasoning continues underneath. The correct value $\frac{1}{8}$ appears only at the final probe. Agreement here is maximal and perfectly uninformative, which is why any deployed rule needs to check answer *shape* before treating agreement as evidence (§5.2).

Case 3: when there is nothing to recover. The minority case is equally worth stating. Problem 240 (gold 116) gives

$$52 \times 2, 104, 52 \times 5, \dots, 154 \times 8, 52, 154 \times 62$$

A three-probe rule fires at token 192 and commits 52; the trajectory itself ends at 154. Both are wrong. Here early stopping destroys nothing, because continued reasoning was not going to arrive at the right answer either—the model settles firmly on 154 and holds it for its final 62 probes. Cases like this are why the harm-to-rescue ratio of §5.3 counts only stops where committing *changes* correctness: on problems the model was never going to solve, a consensus stop costs accuracy nothing and saves tokens, which is the regime consensus early exit was designed for. The difficulty is that agreement does not distinguish this case from Case 1, and Case 1 is the more common one.

G False Consensus: Full Analysis

Section 4.1 summarizes, in four facts, what this appendix establishes in full: in an intervention-free setting, intermediate probe consensus is not a reliable stopping signal.

G.1 Setup

We analyze the frozen main trajectories (Section 3): DeepSeek-R1-Distill-Qwen-7B on all 500 MATH500 problems (Lightman et al., 2024; Hendrycks et al., 2021) at the main 16K-token budget, temperature 0.6, top- p 0.95, across three seeds




Figure 5: **The safe-and-saving corner is empty across the entire rule space.** Dev split, macro-averaged over the 18 environments. Each cell holds the 160 rules that share its (W, s) and differ only in the operational knobs, summarised by the two members a gate interrogates. (a) the cell’s *safest* rule: colour is its accuracy drop and the number beneath is the saving that rule buys. Larger windows and stricter thresholds do drive the drop towards the 1.0 pp cap—but the saving falls with it, reaching -1% at $W=30, s=1.0$. (b) the cell’s *most economical* rule: colour is its net saving and the number beneath is the drop it costs. Substantial saving is available everywhere, at 4–73 pp. A cell would be outlined in green if any of its rules cleared the conservative gate (drop ≤ 1.0 pp and saving $\geq 10\%$); none is.

(1,500 trajectories in total).⁴ The dense probe bank appends the boxed-answer probe suffix every 64 generated tokens and reads off the model’s current answer. The probe is a separate forward pass: the controller only *observes*, never halting or altering the main trajectory. This yields, per trajectory, the full answer sequence a stopping rule would have access to at the budget used throughout the paper.

G.2 Final agreement is reliable; early agreement is not

Agreement measured at the *end* of a trajectory does track correctness. On the 186 trajectories that reach unanimous *cumulative* agreement (all non-empty probes concur), 97.8% are correct—only four whole-trajectory false consensuses; and of the 1,205 whose *final five-probe window* is unanimous, 90.4% are correct. But this end-state agreement is not what an online stopping rule can act on: a rule fires the *first* time it sees agreement, mid-trajectory, long before the answer has settled. The rest of this appendix shows that *early* agreement—the signal a controller must actually use—is far from terminal.

⁴Train/dev problem ids are collected at seeds 42/43/44 and the held-out test ids at seeds 45/46/47; their union covers all 500 problems. This is a descriptive characterization of the frozen trajectories—it selects and tunes nothing—so it does not touch the preregistered sweep or the test-split commitment.

G.3 The cost lands at the stopping decision

We simulate a *naive consecutive-agreement* stop: halt at the first point where three consecutive probes agree, are non-empty, and are *certain* (the model emits the boxed answer with high token-level confidence; formal definition in Appendix A). This is a simplified heuristic, *not* the CertIndex (CoT) rule, which we reproduce from its released implementation in §4.4. It fires on 1,477/1,500 trajectories. The committed answers are correct only 50.5% of the time, whereas letting those same problems run to completion reaches 90.7%—a 40.2 **pp accuracy loss** in exchange for an average saving of 3301 tokens, with 731 trajectories stopped on an outright wrong answer.

G.4 Recovery is common, and early stopping kills it

The loss is driven by *recovery*: continued reasoning overturns the intermediate answer far more often than intuition suggests. Of the 736 trajectories that at some point formed a three-probe consensus *different* from their final answer, 620 (84.2%) ended up correct; and of the 1,148 whose *first* probe was wrong, 89.1% eventually flipped to correct. Neither “first impression” nor “local consensus” is close to terminal.

Counterintuitively, *later* consensus is *less* trustworthy: consensus formed before 512 tokens is 91.0% correct, but consensus that only forms after




Figure 6: **A rule must clear the gate on every split.** The 478 amber rules are exactly those that clear the conservative gate *in-sample on train*; drawing all 478 would be unreadable, so we show their median, interquartile range and full range. Dashed lines mark all three preregistered gates (Conservative, **B**alanced, **T**oken-efficient); a rule must sit below them on the left and above them on the right. On dev, the split the gate is applied on, none of the 478 does. Some recover on test, but a rule selected on dev is never among them: jointly, 0 of 478 clear all three splits, against 3 of 3 for the DEER operating points selected on dev. Macro-averaged over 18 environments per split; train and dev share seeds 42–44 and differ only in problems, while test uses unseen problems and seeds 45–47.

2048 tokens is 71.6% correct. Hard problems both converge slowly and converge to worse answers—exactly the problems on which an aggressive stopping rule does most damage.

and economical. In the searched space, it cannot.

G.5 An error taxonomy

Hand-inspecting stopped-but-wrong cases (a preliminary AI-assisted labeling of 28 cases from an earlier exploratory pass, not the automated counts above) yields *four* non-empty types out of five categories. By count: **(A)** numeric collapse—stable convergence to a wrong number—14/28; **(D)** derivation gaps—dropped roots/cases, unverified solutions—7/28; **(E)** format/option hallucination—a non-multiple-choice problem stably emitting a letter such as “B”—6/28, which is an artifact of the probe mechanism itself rather than a model belief; and **(B)** expression collapse—1/28. Type **(C)** (sign errors) did not occur in this sample. Type E is itself a methodological finding: roughly one in five “false consensus” is a probe formatting artifact, implying that any deployed rule should filter answer *shape*.

Takeaway. These observations—local consensus unreliable, recovery common, late consensus worse—suggest that a good stopping rule must be much more than “agreement.” The rest of the paper asks whether the *best possible* such rule, drawn from a large preregistered space, can be both safe



Figure 7: Selection → generalization across splits (two development models; macro-averaged over environments). Grey: all 3,520 consensus rules. **Orange**: the 478 rules that clear the conservative gate *in-sample on train*. **Green**: the DEER threshold frontier, with the three dev-selected operating points (**Conservative/Balanced/Token-efficient**) as stars. **Purple star**: the oracle upper bound (earliest correct probe).

Out-of-sample generalization (test split) across scale and architecture: dev-selected DEER points (C/B/T) stay in the gate on every model; consensus train candidates do not




Figure 8: Out-of-sample generalization across scale and architecture (all panels test split). The two development models (top) are shown on their held-out test seeds; the two held-out models (bottom) are entirely unseen. Markers as in Figure 7. The dev-selected DEER points stay inside the conservative gate on every model; no consensus rule does.




Figure 9: Out-of-sample generalization across benchmarks (test split, development models). The consensus trade-off and the DEER advantage hold on MATH500 and AMC23; on the small AIME24 set (few problems per environment) the two frontiers are noisier and intertwine. Markers as in Figure 7.