

False Consensus: Why Intermediate Agreement Is Not a Safe Early-Exit Signal for Reasoning LLMs

Anonymous ACL submission

Abstract

Reasoning language models spend most of their tokens *after* the answer is effectively decided, so a popular way to cut inference cost is to periodically probe the partial trajectory for its current answer and stop once probes *agree*. We show this bet is unsafe, and that the failure is a property of the *signal*, not of the stopping threshold. On frozen reasoning trajectories from two models and three math benchmarks we run a *preregistered* sweep of the consensus signal—a window size $\times$ share threshold family plus operational knobs, 3,520 rules—with problem-grouped splits and acceptance gates (a total accuracy-drop cap, a total token-saving floor, and a positive-saving fraction) fixed in advance. The result is distribution-free: **not one of the 3,520 rules clears any gate**; capping the accuracy drop at 1 pp forces net saving to near zero, while demanding a 10% saving forces a multi-point drop, and even large windows only trade the drop for zero saving. The reason is directional, not statistical: intermediate consensus forms *while the answer is still moving*, so a stop destroys a correct-in-the-end answer roughly **35 times more often** than it rescues a wrong one—an *accuracy tax* intrinsic to the trajectory. The failure is specific to *consensus*: swept through the identical pipeline, gates, and token accounting, a method reading a different signal—the model’s own confidence at reasoning-boundary tokens (DEER)—**clears all three gates**, losing as little as 0.3 pp while saving 28%. The consensus frontier reproduces on a held-out test split ($r=0.98$) and on unseen models at $4\times$ the scale and a different architecture, with no consensus rule safe on both splits, while DEER stays admissible; a faithful Certaindex reproduction on identical trajectories collapses by 56–70 pp. Early exit is possible—just not on the consensus signal. We release the protocol, the frozen trajectories, and the full sweep.

1 Introduction

Reasoning language models such as DeepSeek-R1-Distill-Qwen-7B and Qwen3-8B (DeepSeek-AI, 2025; Qwen Team, 2025; OpenAI, 2024) solve hard problems by emitting long chains of thought, often thousands of tokens per query. Much of this computation is spent *after* the model has effectively decided on an answer, which makes *early exit*—halting generation once the answer is settled—an appealing lever for cutting inference cost (Chen et al., 2024; Fu et al., 2024). A widely used family of methods does this by *probing*: at regular intervals the partial trajectory is queried for its current answer (e.g., by appending a short “Final Answer:” suffix), and generation is halted once several probes *agree* (Fu et al., 2024). The bet is that *intermediate consensus signals termination*—once the model repeatedly says the answer is x , further reasoning is unlikely to change it. This paper asks a single question: **is intermediate agreement a safe signal on which to stop?**¹

Our answer is **no**, and the reason is not a badly tuned threshold—it is the signal itself. Intermediate consensus forms *while the answer is still moving*: on a frozen replay, stopping on consensus destroys a correct-in-the-end answer roughly **35 times more often** than it rescues a wrong one (Section 5). This is an *accuracy tax* that is a property of the reasoning trajectory, not of how one probes it, so no threshold on the agreement signal can escape it. We make the claim hard to dismiss as “you did not tune it well” by searching the space exhaustively under commitment: on frozen trajectories from two models and three benchmarks we run a *preregistered* sweep of the consensus signal—a

¹We use “false consensus” for the phenomenon that intermediate probe agreement is non-terminal; the term is unrelated to the social-psychology “false consensus effect” (Ross et al., 1977). We study answer-level early exit (stopping a chain of thought), not token- or layer-level early exit (Graves, 2016; Schuster et al., 2022).

window size $\times$ share threshold family plus operational knobs, 3,520 rules—with problem-grouped splits and acceptance gates fixed in advance (Sections 3–4). The outcome does not depend on any single fragile estimate: **not one of the 3,520 rules is both safe and saving—none clears any gate**, and even large windows only trade the drop away for zero saving. To show this is the *signal* and not early exit, we sweep one non-consensus method—DEER, which reads the model’s confidence at a reasoning boundary—through the identical pipeline and gates: it clears all three, losing as little as 0.3 pp while saving 28%. The consensus frontier reproduces on a held-out test split ($r=0.98$) and on unseen models at $4\times$ the scale and a different architecture (Section 4), and a faithful Certainindex reproduction on the same trajectories collapses by 56–70 pp. Figure 1 is the whole paper in one picture: the safe-and-saving corner is empty for every consensus rule, while the one method that clears the gates is the one that does *not* read probe agreement.

Contribution. Our contribution is a rigorous, mechanism-backed demonstration that *intermediate answer-consensus is an unsafe early-exit signal for reasoning LLMs*, established so that it cannot be waved away as an untuned heuristic. Three components support this one claim:

- **A preregistered, probe-independent test** (frozen trajectories, offline re-probing, a two-parameter consensus family enumerated to 3,520 rules, and selection gates fixed in advance) under which *no* consensus rule is both safe and token-saving—and which we release in full (Sections 3–4).
- **A mechanism** that explains *why*: a directional “continued-reasoning-corrects” effect ($\approx 35:1$) makes the accuracy cost of stopping intrinsic to the trajectory and independent of probe density, so the negative result generalizes beyond the searched schema (Section 5).
- **Out-of-distribution and cross-signal confirmation:** the consensus frontier reproduces on a held-out split and two unseen models, a published Certainindex reproduction collapses, and a non-consensus signal (DEER) swept under the identical protocol *clears the same gates* on the same trajectories—localizing the failure to consensus and indicating where a safe signal must come from (Sections 4–5).

2 Related Work

Probe-based / consensus early exit. The methods we scrutinize periodically query a partial reasoning trajectory for its current answer and halt on agreement. Dynasor / Certainindex (Fu et al., 2024) formalize this with a “certainty index” over recent probe answers and use it to terminate early when serving reasoning programs, one of several responses to the *overthinking* of reasoning models (Chen et al., 2024; Yang et al., 2025). We adopt the same probe mechanism (a short boxed-answer suffix) and notion of probe agreement, and ask how far *any* rule in this family can be pushed before it becomes unsafe.

Early-stopping self-consistency and self-correction. Probe-based early exit is the online, *single-trajectory* analogue of early-stopping self-consistency (Wang et al., 2023; Aggarwal et al., 2023; Li et al., 2024): it votes across snapshots of one in-progress trajectory rather than across finished samples. That difference is the crux—snapshots of an unfinished trajectory are not exchangeable draws from a stable answer distribution, because the trajectory is still correcting itself. This intrinsic self-correction differs from the prompted, post-hoc self-correction that Huang et al. (2024) find unreliable: we measure answer evolution *within* one uninterrupted trajectory, and find (Section 4) that a model’s mid-trajectory answer is far from final.

Alternative termination signals. Rather than probe agreement, a stop can read the model’s own forward-looking state: verifier- and process-reward-guided decoding (Lightman et al., 2024; Wang et al., 2024) score partial reasoning, and adaptive-computation methods learn a halting signal at the token or layer level (Graves, 2016; Schuster et al., 2022). DEER (Yang et al., 2025) stops on the model’s confidence in a trial answer at reasoning boundaries. Our results (Section 5) indicate that this class—forward-looking signals about whether the trajectory will still change—is the direction a safe stop must take, in contrast to the backward-looking agreement signal we show fails. Test-time scaling trades compute for accuracy via repeated sampling and search (Snell et al., 2024; Brown et al., 2024); early exit is its dual, spending less without losing accuracy.

Preregistration. To keep an exhaustive rule search from silently overfitting to a favorable op-




Figure 1: **The safe-and-saving corner is empty for consensus; a different signal fills it.** Each grey point is one of the 3,520 preregistered consensus rules on the development set (a two-parameter window size $\times$ share threshold family plus operational knobs); the red step is their dev Pareto boundary. The top-left corner—small accuracy drop *and* large token saving—contains no consensus rule: the boundary only climbs as the drop grows, and *none* of the three shaded gate regions is entered. **DEER** (Yang et al., 2025), swept over its confidence threshold through the same pipeline and accounting but reading a non-consensus signal, traces the green frontier that sits above the consensus boundary throughout the safe region and clears every gate. The failure is a property of the consensus signal, not of early exit.

erating point, we fix the rule schema, splits, and acceptance gates before looking at held-out data (Nosek et al., 2025). To our knowledge this is among the first preregistered evaluations of early-exit rules for reasoning models; the protocol (Section 3) is what makes a *negative* result interpretable.

3 Method

Section 2 refuted one heuristic. To ask whether *any* probe-based rule can succeed, we need (i) an evaluation that a rule cannot game by changing the reasoning it observes, and (ii) a selection procedure that cannot silently overfit to a lucky operating point. We preregister both.

3.1 Probe-independent, frozen trajectories

For each problem we generate exactly *one* main trajectory in a single request (caps: 16K tokens for MATH500 and AMC23, 32K for AIME24) and freeze it. All probing is then performed *offline* against fixed text prefixes of that frozen trajectory. Because probes never re-enter the decode, changing a probe schedule—or a stopping rule—cannot change the underlying reasoning; every rule is scored against the same trajectories. We build two offline probe banks: a *dense* bank that re-probes every 64 tokens with simple@32 (a 32-sample boxed-answer probe), and an *adaptive* bank that additionally fires event-triggered probes at entropy drops and at conclusion/reflection/answer markers.

3.2 Token accounting

For a rule that stops a problem at token position s after consuming probe output p , we charge total decode tokens $T = s + p$ and compare against the frozen baseline B (the full trajectory, capped at budget). *Net* savings is $(B - T)/B$ and *gross* savings is $(B - s)/B$; their gap is the probe cost. The accuracy drop compares the correctness of the committed answer at s against that of the frozen final answer (it is not an agreement rate between them). Two quantities therefore move independently, and we track both: *where* the rule stops (setting the accuracy drop and the main-token saving s) and *how much it probes* to get there (adding p). Section 5 shows that only the first is intrinsic.

3.3 The consensus signal: window size and share threshold

The consensus signal any probe-based rule can read reduces to two hyperparameters: a **window size** W (how many of the most recent probes to inspect) and a **share threshold** s (what fraction of that window must agree on a single answer before a stop is allowed). This pair spans the family exactly: $W=1$ is the “latest-probe” rule, $s=1.0$ is strict unanimity over the last W probes (the Certainindex-style stop), and an entropy-over-the-window criterion is monotone in the same agreement and adds no reachable operating point. We sweep $W \in \{1, 3, 5, 8, 12, 16, 24, 30\}$ and $s \in \{0.6, 0.8, 1.0\}$, deliberately including *large* windows so a rule can demand a long, stable agreement before committing.

Around this core we vary only the operational knobs that govern *when* and *how* the signal is read: **probe schedule** (fixed interval $\in \{64, 128, 256, 512\}$ tokens, or event-triggered probing at conclusion/entropy/reflection markers); **maturity** (a minimum-token floor $\in \{0, 512, 1024, 2048, 4096\}$ before any stop is allowed); **validity** (an answer-shape filter, rejecting empty/letter answers or only empties); and **certain** (optionally requiring the supporting probes to be flagged certain). Enumerating this preregistered grid over a fixed-schedule and an event-triggered family yields 3,520 candidate consensus rules; formal definitions are in Appendix A.

A better signal, swept identically. To test whether the failure is early exit *itself* or the consensus *signal*, we sweep one non-consensus method through the same pipeline, gates, and token ac-

counting. **DEER** (Yang et al., 2025) stops on the model’s own confidence in a trial answer at a reasoning boundary; its only hyperparameter is a confidence threshold, which we sweep over 14 values. We use the stronger *trial-answer-submit* variant (commit the trial answer directly, no separate readout) rather than the faithful readout reproduction—our aim here is not to reproduce prior work but to ask whether *any* rule can clear the gates. Placing DEER and consensus on one frontier under identical accounting is what makes the comparison a controlled test of the signal.

3.4 Models, benchmarks, splits

Development uses two models—DeepSeek-R1-Distill-Qwen-7B and Qwen3-8B—on three benchmarks (MATH500, AMC23, AIME24), each at seeds 42/43/44: 18 model $\times$ benchmark $\times$ seed environments. Problems are partitioned *by problem id* into 60/20/20 train/dev/test splits (MATH500 300/100/100, AMC23 24/8/8, AIME24 18/6/6), so no problem leaks across splits. Confirmation reserves seeds 45/46/47 plus two unseen models—DeepSeek-R1-Distill-Llama-8B and Distill-Qwen-32B—all evaluated on the *test* split only.

3.5 Preregistered operating points and gates

We fix three operating points and their acceptance gates *before* looking at held-out data (Table 1). Each gate is applied in order: first the **total accuracy drop** (macro-averaged over the 18 environments) must stay under a cap; then the **total net token saving** (macro-averaged) must clear a floor; and finally a **positive-saving fraction** (psf, the fraction of environments with strictly positive net saving) floor must be met. A rule is accepted at an operating point only if it satisfies all three on train+dev. The saving floor is what distinguishes a genuinely useful early exit from a rule that stays accurate only by almost never stopping.

3.6 Metric resolution

The total accuracy drop is a macro-average over 18 environments, three of them per benchmark $\times$ model, so the small competition sets do not dominate: one AIME24 problem is 16.7 pp of a 6-problem dev split but contributes only its share of one of the 18 cells. The gates are read on the dev split (the selection target) with train reported alongside; because the central finding is distribution-free (*no* consensus rule clears any gate,

Operating point	Max total acc. drop	Min total saving	psf floor
conservative	1.0 pp	10%	0.80
balanced	2.0 pp	20%	0.80
token_efficient	3.5 pp	30%	0.70

Table 1: Preregistered acceptance gates. A rule must keep the macro-averaged accuracy drop under “max total acc. drop,” deliver at least “min total saving” net tokens, and post positive net saving on at least a psf fraction of environments. Gates are fixed in advance and *not* relaxed post hoc.

by a wide margin, Section 4), it does not rest on a single fragile cell. Held-out confirmation on the test split and two unseen models provides the complementary out-of-distribution check.

3.7 Preregistration commitments

Three commitments make the result interpretable. **(C1)** The test split and both confirmation models are never read during rule selection. **(C2)** The conservative gate in Table 1 is not relaxed after seeing results; if no rule passes, that is reported as the finding rather than engineered away. The balanced and token_efficient points are reported only as clearly-labeled post-hoc analysis, never as a selected operating point. **(C3)** The frozen rule set is fixed on train+dev and its hash recorded before any test/confirmation evaluation. Section 4 reports what happened when we applied C2.

4 Experiments

We first confirm the phenomenon the paper turns on—intermediate consensus is non-terminal (§4.1)—then report the preregistered sweep and its central negative result: *no* consensus rule is both safe and saving (§4.2). We then show the failure is specific to the *signal*, not to early exit, by sweeping a non-consensus method (DEER) through the identical pipeline and gates and finding that it *does* clear them (§4.3); a faithful Certainindex reproduction on the same trajectories collapses (§4.4). Finally we confirm the frontier out of distribution (§4.5).

4.1 False consensus is real

We log a dense probe stream for every MATH500 problem under DeepSeek-R1-Distill-Qwen-7B, never intervening—the probe is a separate forward pass that only observes the main trajectory (full setup in Appendix E)—and read off four facts.

(1) Agreement is not correctness. *Cumulative* agreement (all non-empty probes so far concur) is

nearly trustworthy (98.9% correct on the 87 problems reaching it), but that is not what an online rule sees. The operative signal, *window* agreement (the last five probes concurring), is markedly less reliable: of the 338 problems reaching a unanimous last-five window, only 93.5% are correct—a 6.5% false-consensus rate. **(2) Recovery is common.** Of the 375 problems whose *first* probe is wrong, 76.3% eventually flip to correct; of the 145 that formed a three-probe consensus *different* from their final answer, 65.5% ended up correct. **(3) Later consensus is worse.** Consensus formed before 512 tokens is 87.4% correct, but consensus forming only after 2048 tokens is 58.1% correct—hard problems converge slowly *and* to worse answers. **(4) The cost lands at the stop.** A naive stop (first three consecutive agreeing, certain, non-empty probes) fires on 416/500 problems and commits to answers correct only 69.2% of the time, whereas letting those same problems run to completion reaches 85.6%: a 16.4 pp loss.²

A good stopping rule must therefore be much more than “agreement.” The rest of this section asks whether the *best possible* such rule, drawn from a large preregistered space, can be both safe and economical.

4.2 The safe-and-saving corner is empty for consensus

We replay all 3,520 consensus rules on the frozen trajectories and aggregate to per-environment metrics—126,720 rows = 3,520 rules $\times$ 18 environments $\times$ 2 splits (train, dev)—then, per rule, its total (macro) accuracy drop and net token saving (accounting of Section 3.2). Applying the preregistered gates yields the paper’s central negative result.

The outcome is distribution-free: across all 126,720 rule–environment evaluations accuracy strictly *falls* in 64.9% and strictly *rises* in only 10.5% (unchanged in 24.6%), for a mean drop of 13.0 pp; over rules the total dev drop has median 13.2 pp and a minimum of 0.81 pp. **Not one of the 3,520 rules clears any gate.** The reason is the shape of the frontier (Figure 1): staying accurate and saving tokens are in direct tension. Only 40 rules keep the total drop at or below the conser-

²This illustration uses a tight 3072-token budget so truncation is visible on a single A100 and a simplified heuristic, *not* the Certainindex algorithm (§4.4). The main study below runs at the 16K/32K budgets used throughout, where the same effect appears as the directional accuracy tax of Section 5.

Method	Acc.	Δ acc (macro)	Net saving
Full generation	82.5%	—	0%
Consensus (drop $\leq$ 1.0)	81.6%	−0.9 pp	+0.2%
Consensus (save $\geq$ 10%)	79.9%	−2.7 pp	+10.7%
DEER (conservative)	82.5%	−0.33 pp	+28.2%
DEER (balanced)	81.9%	−1.03 pp	+29.6%
DEER (token_eff.)	80.1%	−2.75 pp	+31.9%

Table 2: Dev operating points (macro over 18 environments). No consensus rule is simultaneously safe (≤ 1.0 pp) and saving ($\geq 10\%$): capping the drop forces saving to near zero, and demanding saving forces a large drop. DEER, reading a non-consensus signal, clears all three gates on the same trajectories under the same accounting.

vative 1.0 pp cap, and the most any of them saves is 0.2%—far under the 10% floor; conversely, the first rule to save 10% already costs 2.66 pp, saving 20% costs 6.17 pp, and saving 30% costs 11.8 pp. Crucially, the sweep now includes *large* windows (up to $W=30$): these do buy low accuracy drop, but only by stopping so late that net saving collapses to zero. No setting of the consensus signal escapes the trade-off; the safe-and-saving corner is empty.

4.3 Early exit is possible—with a different signal

The same gates are cleared, comfortably, by a rule that does *not* read probe agreement. Sweeping DEER’s confidence threshold through the identical pipeline and token accounting (Section 3.3), its trial-answer-submit variant passes all three operating points on dev: the conservative point loses only 0.33 pp while saving 28.2%, the balanced point 1.03 pp at 29.6%, and the token-efficient point 2.75 pp at 31.9%; at a stricter threshold it is accuracy-neutral (−0.06 pp) while still saving 20.8%. In Figure 1 the DEER frontier sits *above* the consensus frontier everywhere in the safe region and enters every gate box that no consensus rule reaches. Early exit is therefore not impossible; the consensus *signal* is what cannot make it safe (Table 2).

4.4 Consensus in the wild: a faithful Certaindex reproduction collapses

The sweep shows no *tuned* consensus rule is safe; a published consensus method shows the same failure at its shipped operating point. We reproduce three prior stoppers in a common harness—

the same frozen trajectories, the same simple@32 probe bank, the same token accounting. **CertaIndex** (Fu et al., 2024) is the faithful share-threshold stop (the direct descendant of the consensus family); **DEER** (Yang et al., 2025) here is the *faithful* readout reproduction (the stronger trial-answer-submit variant is what we swept in §4.3); **TJE** (Anonymous, 2026) is a token-level stop. Table 3 reports the reproductions (not the original papers’ end-to-end runs).

CertaIndex, at 99.8% stop rate and 70 pp lost, is the accuracy tax paid in full—the shipped consensus method landing exactly where the swept consensus frontier says it must. TJE is unsafe on the weaker model. Even the *faithful* DEER readout, a non-consensus signal, already avoids collapse; the stronger trial-answer-submit variant we swept turns that into a gate-clearing frontier (§4.3). That only the non-consensus signal survives is the crux we analyze in Section 5.

4.5 The frontier reproduces out of distribution

We read the test split once (commitment C3). Because the development result is that *no* consensus rule clears any gate, there is nothing to confirm by re-selection; we ask whether the *frontier* reproduces, ranking rules on dev and measuring on test.

Held-out test split. Each consensus rule’s total drop on test tracks its dev value closely ($r = 0.98$). The split-invariant claim is the empty *joint* gate: **no consensus rule clears the conservative gate on both splits**—0 on dev, and the 444 rules admissible on *test alone* are exactly the in-sample winners the held-out split exists to reject (all fail on dev). DEER, by contrast, remains admissible across splits: 3 of its thresholds clear the conservative gate on *both* dev and test (and 5 clear balanced on both), so the cross-signal contrast is not a dev artifact.

Unseen scale (4 $\times$) and architecture. Both unseen models are evaluated on the test split at three seeds (45/46/47; nine environments each), so their gates are as well resolved as development’s. On Distill-Qwen-32B, never used in development, the consensus frontier reproduces ($r = 0.97$ against dev drop); on the held-out architecture and family DeepSeek-R1-Distill-Llama-8B it reproduces at $r = 0.94$.³ The *conservative* gate stays empty

³The Llama-8B trajectories were re-collected after correcting a chat-template defect (a missing model-specific

Method	Model	Accuracy	Δacc	Fair token saving	Stop rate	Signal
Full generation	Qwen3-8B	85.4%	—	0%	0%	—
Full generation	DeepSeek-R1-Distill-Qwen-7B	79.8%	—	0%	0%	—
CertaIndex	Qwen3-8B	15.3%	−70.1 pp	90.1%	99.8%	consensus
CertaIndex	DeepSeek-R1-Distill-Qwen-7B	23.9%	−55.9 pp	76.7%	98.7%	consensus
TJE	Qwen3-8B	85.0%	−0.4 pp	2.0%	22.5%	token-level
TJE	DeepSeek-R1-Distill-Qwen-7B	60.7%	−19.1 pp	65.0%	93.4%	token-level
DEER	Qwen3-8B	86.2%	+0.8 pp	16.3%	41.4%	boundary conf.
DEER	DeepSeek-R1-Distill-Qwen-7B	74.9%	−4.8 pp	20.2%	56.1%	boundary conf.

Table 3: Prior early-stoppers reproduced on our frozen trajectories (dev, macro over three benchmarks; probe tokens charged). CertaIndex—the consensus descendant—**collapses** (−56 to −70 pp) while saving 77–90% of tokens: it fires on almost every problem, very early, exactly the false-consensus trap of §4.1. **DEER**, which reads a *different* signal (boundary confidence), is the only method that stays near full accuracy—+0.8 pp on Qwen3-8B—while still saving tokens. We interpret this contrast in Section 5.

on every model—dev, 32B, and Llama all admit 0 rules (the best rule under a 1.0 pp drop saves 0.6% on 32B and 9.3% on Llama). At the looser gates a scale effect appears: on the $4\times$ larger 32B a few consensus rules become admissible in-sample (4 at balanced, 6 at token-efficient), reflecting that a bigger model’s intermediate answers stabilize somewhat earlier—but these are not the dev-selected rules (dev admits none), and on the same-scale Llama architecture the gates stay empty (0/0/0). Consensus thus yields no rule that is safe-and-saving where it is selected, on any model. DEER, swept identically on these unseen models, **clears the gates on both**: on Qwen-32B it is accuracy-neutral at the conservative point (−0.24 pp while saving 32.4%, $\tau=0.97$) and on Llama-8B it loses 0.67 pp at 26.7% ($\tau=0.99$), passing all three gates. Consensus is unsafe, and DEER safe-and-saving, across all four axes we vary—benchmark, seed, $4\times$ scale, and architecture/family.

5 Analysis

Section 4 is an existence proof over 3,520 rules, two seen and two unseen models, and both splits. This section explains *why* it holds and why it generalizes beyond the exact rules we searched: the accuracy cost of stopping on consensus is a directional property of the reasoning trajectory itself, not of a threshold or of probe density.

5.1 Stopping destroys more than it saves: a 35:1 effect

Consider only the stops where committing early *changes* correctness relative to the trajectory’s

beginning-of-sequence token that degraded its generations); the corrected model scores $\sim 88\%$ on MATH500.

own final answer, and count two outcomes: (final-correct, stop-wrong)—a correct recovery destroyed—versus (final-wrong, stop-correct)—a wrong answer luckily banked. Pure sampling noise predicts a $\approx 1:1$ ratio; a genuine “continued reasoning corrects” effect predicts $\gg 1$. On the frozen replay this ratio is **35:1** for the naive consensus stopper (1055 destroyed vs. 30 banked; 34.3:1 and 35.9:1 on DeepSeek-R1-Distill-Qwen-7B and Qwen3-8B separately) and 15–18:1 for the more conservative variants—all far above 1:1. Stopping on consensus removes a correct-in-the-end answer about 35 times more often than it rescues a wrong one. This is the recovery phenomenon of §4.1 (65.5%, 76.3%), now at the main 16K/32K budgets and as a signed effect: the loss is real and directional, not sampling variance.

5.2 The accuracy tax is intrinsic to the trajectory

Why is the ratio structural? For a rule stopping at token position s , the accuracy drop depends only on s : committing at s forgoes whatever correction the trajectory would have made after it. It is a property of the reasoning, not of the probe, so it survives *any* probing scheme. This is what never reaches zero in the sweep. The probe grid (intervals 64/128/256/512 tokens) makes essentially every stop position reachable and the observed floor is the minimum over these schedules; a sparser schedule can only stop later or at a subset of positions, never at a safer one the grid missed. Hence no threshold on the consensus signal—in or beyond our schema—escapes the accuracy tax; the searched space is a lower bound on the cost, not a lucky miss.

Consensus point	Gross saving ($B - s$)/ B	Net saving ($B - s - p$)/ B
save $\geq$ 10% (drop 2.7 pp)	+14.7%	+10.7%
save $\geq$ 20% (drop 6.2 pp)	+27.0%	+20.2%

Table 4: Gross vs. net savings for two consensus frontier points; the gap (~ 4 –7 pp) is the probe tax. It shifts the *savings* axis but never rescues the accuracy drop—which is the intrinsic tax that keeps the corner empty.

5.3 Ruling out the “it is just dense probing” objection

The natural rebuttal is that the negative result is an artifact of charging dense re-probing. It is not, and separating two costs shows why. Besides the accuracy tax, a rule pays a *probe tax*: the probe output p (Section 3.2). Under dense re-probing a late-stopping large-window rule pays p over a long prefix, opening a gap between its *gross* savings and its (smaller) *net* savings (Table 4). Unlike the accuracy tax, the probe tax scales with probe density and is reducible (sparser or shorter probes, KV-cache reuse). The two are separable: the *savings* axis is partly a probe-tax artifact and is addressable, but the *accuracy* floor is probe-independent and is what makes the corner empty—and it is the accuracy floor, not the probe tax, that DEER escapes.

5.4 It is the signal, not early exit

The joint sweep of §4.3 localizes the cause. Every method that stops on *probe agreement*—the swept consensus family and CertIndex alike—either fails a gate or collapses outright; CertIndex’s 70 pp loss at 99.8% stop rate is the accuracy tax paid in full. The method that clears the gates, DEER, does not read probe agreement at all: it stops on the model’s own confidence in a trial answer at a reasoning boundary, and in Figure 1 its threshold frontier sits *above* the consensus Pareto boundary throughout the safe region. Because DEER is swept through the identical pipeline, gates, and token accounting, the contrast is a controlled test: the same early-exit machinery succeeds on a forward-looking signal and fails on the consensus signal. The failure is a property of the *signal*, not of early exit.

5.5 Implication: change the signal, not the threshold

Much prior effort tunes consensus-based stops (Fu et al., 2024; Aggarwal et al., 2023; Li et al., 2024); our frontier shows the ceiling of that approach is

low, and the mechanism above says no amount of tuning raises it. The constructive implication is that a safe signal should be about whether the *trajectory itself* will still change, not about the agreement of backward-looking probes. Our joint sweep already demonstrates this: the boundary-confidence signal (DEER) clears the same gates the whole consensus family fails (§4.3). As a further step toward an online controller, we sketch in Appendix F a fast-commit/verify variant built on the same signal; we stress it is exploratory (three seeds, seed-sensitive, outside the preregistration) and report it only as a design direction, not a finished method. Verifier- and process-reward-based termination (Lightman et al., 2024; Wang et al., 2024) and learned halting (Graves, 2016; Schuster et al., 2022) point the same way. Turning any of these into a *confirmed* method—with its own train/dev→test validation—is the natural next paper; our contribution is to establish rigorously which signal not to build it on. Two caveats bound the claim: it concerns *consensus-based* rules within the searched schema and dense-probe regime, and it does not say early exit is impossible—only that this signal cannot make it safe.

6 Conclusion

We asked whether intermediate answer-consensus is a safe signal on which to early-exit a reasoning LLM. It is not, and not because the rule is untuned: intermediate consensus forms while the answer is still moving, so stopping on it destroys a correct-in-the-end answer about 35 times more often than it rescues a wrong one. This accuracy tax is a property of the trajectory rather than the probe, and it shows up as an empty safe-and-saving corner across a preregistered space of 3,520 rules—not one of which clears any gate—reproduced on a held-out split and two unseen models, and paid in full by a faithful CertIndex reproduction that collapses by 56–70 pp. That the corner is empty for consensus but not for early exit is the point: swept through the identical protocol, a method reading the model’s own boundary confidence (DEER) clears all three gates on the very same trajectories, so the failure is specific to consensus. The takeaway for anyone building safe early exit is therefore about the choice of signal: read the model’s own forward-looking state, not the agreement of its backward-looking probes.

Limitations

Scope of the negative result. The central result—no consensus rule clears any gate—is established on 18 development environments (two models, three benchmarks, seeds 42/43/44) and confirmed on the held-out test split and two unseen models (Section 4). It is a statement about the space searched: one boxed-answer probe suffix (simple@32) and one preregistered consensus schema (3,520 rules over a window-size $\times$ share-threshold family plus operational knobs). Other probe wordings and signals outside this schema are not searched; the mechanism (Section 5) is our argument that consensus variants would not help, but it is an argument, not an exhaustive proof.

DEER as the positive control, not a benchmarked SOTA. We sweep DEER’s trial-answer-submit variant to establish that *some* signal clears the gates the consensus family fails (on the development models, the held-out test split, and both unseen models); we do not claim it is the best possible early-exit method. Its role is a controlled positive control for the signal, not a leaderboard entry.

Held-out architecture was re-collected. The held-out architecture model (Llama-8B) was re-collected after we found and fixed a chat-template defect (a missing model-specific beginning-of-sequence token) that had degraded its earlier generations; the corrected model scores $\sim 88\%$ on MATH500, and the confirmation uses the corrected trajectories.

Small held-out sets on hard benchmarks. AIME24 and AMC23 are small (dev splits of 6 and 8 problems per seed), and the unseen models contribute a single seed, so those per-environment cells are noisy. The total drop macro-averages over the 18 development environments (three seeds per benchmark $\times$ model), and we read the single-seed unseen models as frontier-reproduction evidence rather than as independent gate tests (Section 3.6).

Probe-tax dependence of the savings axis. Net-savings numbers charge dense re-probing (every 64 tokens, 32-token probes); a sparser or KV-cache-reusing schedule would change the *savings* axis and could move some rules to positive net savings. We separate gross from net savings (Section 5) precisely so the probe-tax-dependent part is distinguishable; the accuracy-tax result does not depend

on the schedule.

The constructive signal is exploratory. The boundary-confidence controller of Appendix F spans three dev seeds but is seed-sensitive (macro Δ acc per-seed range $[-6.1, +4.2]$ pp; mean -0.75 pp), model-dependent, not part of the pre-registered sweep, and has no test-split confirmation. It should be read as evidence that a non-consensus signal escapes the accuracy tax, not as a finished method.

Domain. All benchmarks are competition mathematics with checkable final answers. Whether false consensus and the accuracy tax transfer to open-ended reasoning, code, or agentic tasks is out of scope.

References

- Pranjal Aggarwal, Aman Madaan, Yiming Yang, and Mausam. 2023. Let’s sample step by step: Adaptive-consistency for efficient reasoning and coding with LLMs. In *Proceedings of EMNLP*.
- Anonymous. 2026. [Token-level justification for early-exit in reasoning](#). In *Findings of EACL*.
- Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré, and Azalia Mirhoseini. 2024. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv preprint arXiv:2407.21787*.
- Xingyu Chen, Jiahao Xu, Tian Liang, Zhiwei He, Jianhui Pang, Dian Yu, Linfeng Song, Qiuzhi Liu, Mengfei Zhou, Zhuosheng Zhang, Rui Wang, Zhaopeng Tu, Haitao Mi, and Dong Yu. 2024. Do not think that much for $2+3=?$ on the overthinking of o1-like LLMs. *arXiv preprint arXiv:2412.21187*.
- DeepSeek-AI. 2025. [DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning](#). *arXiv preprint arXiv:2501.12948*.
- Yichao Fu, Junda Chen, Siqi Zhu, Zheyu Fu, Zhongdongming Dai, Aurick Qiao, and Hao Zhang. 2024. [Efficiently serving LLM reasoning programs with Certaindex](#). *arXiv preprint arXiv:2412.20993*.
- Alex Graves. 2016. Adaptive computation time for recurrent neural networks. *arXiv preprint arXiv:1603.08983*.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. Measuring mathematical problem solving with the MATH dataset. In *NeurIPS Datasets and Benchmarks Track*.

Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Wei Yu, Xinying Song, and Denny Zhou. 2024. Large language models cannot self-correct reasoning yet. In *International Conference on Learning Representations (ICLR)*.

Yiwei Li, Peiwen Yuan, Shaoxiong Feng, Boyuan Pan, Xinglin Wang, Bin Sun, Heda Wang, and Kan Li. 2024. Escape sky-high cost: Early-stopping self-consistency for multi-step reasoning. In *International Conference on Learning Representations (ICLR)*.

Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. 2024. Let’s verify step by step. In *International Conference on Learning Representations (ICLR)*.

Brian A. Nosek et al. 2025. Preregistration in machine learning research. *arXiv preprint*.

OpenAI. 2024. Learning to reason with LLMs. <https://openai.com/index/learning-to-reason-with-llms/>.

Qwen Team. 2025. *Qwen3 technical report*. *arXiv preprint arXiv:2505.09388*.

Lee Ross, David Greene, and Pamela House. 1977. The “false consensus effect”: An egocentric bias in social perception and attribution processes. *Journal of Experimental Social Psychology*, 13(3):279–301.

Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q. Tran, Yi Tay, and Donald Metzler. 2022. Confident adaptive language modeling. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2024. Scaling LLM test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*.

Peiyi Wang, Lei Li, Zhihong Shao, Runxin Xu, Damai Dai, Yifei Li, Deli Chen, Yu Wu, and Zhifang Sui. 2024. Math-shepherd: Verify and reinforce LLMs step-by-step without human annotations. In *Proceedings of ACL*.

Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations (ICLR)*.

Chenxu Yang, Qingyi Si, Yongjie Duan, Zheliang Zhu, Chenyu Zhu, Zheng Lin, Li Cao, and Weiping Wang. 2025. *Dynamic early exit in reasoning models*. *arXiv preprint arXiv:2504.15895*.

A Rule Schema Details

The consensus schema of Section 3 centers on two hyperparameters and a few operational knobs, enumerating to 3,520 candidate rules over a fixed-schedule and an event-triggered family. Each rule fixes:

1. **Window size** $W \in \{1, 3, 5, 8, 12, 16, 24, 30\}$: the number of most recent probes inspected. $W=1$ is the latest-probe rule.
2. **Share threshold** $s \in \{0.6, 0.8, 1.0\}$: the fraction of the last W probes that must agree on one answer (share measured over the full window, so an empty or dissenting probe counts against consensus); $s=1.0$ is strict unanimity. An entropy-over-the-window criterion is monotone in the same agreement and adds no reachable operating point, so it is not swept separately.
3. **Probe schedule**: fixed interval $\in \{64, 128, 256, 512\}$ tokens, or adaptive_event with entropy-drop and marker triggers and a fallback interval.
4. **Maturity**: a minimum-token floor $\in \{0, 512, 1024, 2048, 4096\}$ before any stop is allowed.
5. **Certainty**: optionally require the supporting probes flagged certain (minimum fraction).
6. **Validity**: answer-shape filter (reject empty / single-letter answers to suppress the Type-E probe artifact of Appendix E), or accept any non-empty answer.

Certainty. A probe is *certain* when its boxed answer is emitted with high token-level confidence: concretely, the mean probability of the answer tokens under the probe forward pass exceeds a fixed threshold (the Certainindex certainty flag). For the simple@32 bank this is aggregated over the 32 samples as the fraction agreeing on the modal answer. Dimension 5 gates a stop on the fraction of the supporting probes that are certain.

Entropy triggers. The adaptive_event schedule fires an extra probe when the token-level predictive entropy of the main decode drops by at least ΔH relative to a trailing baseline (an “the

model just committed” signal), and at conclusion/reflection/answer marker tokens (e.g., “therefore”, “the answer is”), subject to a minimum inter-probe gap and a fallback interval.

B Preregistration Text

The protocol version, split manifest hash, candidate-rule generation script, and the three operating-point gates (Table 1) were fixed before any held-out evaluation. Commitments C1–C3 (Section 3) were recorded with the protocol. When the conservative gate proved empty, we followed C2 and report the empty gate as the finding rather than relaxing it; the balanced/token_efficient points are reported only as clearly-labeled post-hoc analysis and would be re-preregistered if used for selection.

C Reproducibility

Frozen trajectories, offline probe banks, the sweep archive (126,720 dev metric rows plus the DEER threshold sweep), and the aggregation scripts are released. Tables 5–6 were reproduced to the reported precision directly from the released sweep archive. Grading uses a robust answer-equivalence check that tries the raw and normalized reference forms and both string and math-expression equality; per-problem correctness for the frozen baseline is recomputed at replay time rather than trusted from collection.

D Window-Size and DEER Frontiers

Table 5 traces the consensus trade-off directly along the window axis: as the window W grows, the minimum attainable accuracy drop falls monotonically (from 22 pp at $W=1$ to 0.8 pp at $W=30$), but the maximum net saving falls with it (from 98% to 39%)—the low-drop region is reachable only at large windows, where saving has already collapsed. This is why no single (W, s) setting lands in the safe-and-saving corner. Table 6 gives the DEER threshold frontier that does reach it.

E False Consensus: Full Analysis

Section 2 summarizes the four facts this appendix establishes in full, in an intervention-free setting: that intermediate probe consensus is not a reliable stopping signal.

Window W	Min total drop (pp)	Max net saving
1	22.25	98.1%
3	11.76	93.8%
5	8.01	87.9%
8	4.91	72.9%
12	1.85	62.0%
16	1.74	55.1%
24	1.69	43.2%
30	0.81	38.8%

Table 5: Consensus frontier along the window axis on dev (over all share thresholds and operational knobs at each W). Larger windows buy a smaller minimum accuracy drop but only by stopping later, so the maximum net saving falls in lockstep; the two never meet in the gate region.

Threshold τ	Total drop (pp)	Net saving
0.9999	−0.06	20.8%
0.999	0.06	25.4%
0.995	0.33	28.2%
0.99	1.03	29.6%
0.97	2.75	31.9%
0.95	5.87	34.8%
0.90	11.43	42.6%

Table 6: DEER trial-answer-submit threshold frontier on dev (macro over 18 environments, same token accounting). Unlike consensus, low drop coincides with substantial saving—the conservative/balanced/token-efficient gates are all cleared around $\tau \in [0.97, 0.995]$.

E.1 Setup

We run DeepSeek-R1-Distill-Qwen-7B (vLLM, one A100-80G) on all 500 MATH500 problems (Lightman et al., 2024; Hendrycks et al., 2021) with a decode budget of 3072 tokens, temperature 0.6, top- p 0.95, and seed 42. Every 128 generated tokens (at most 24 probes per problem) we append the boxed-answer probe suffix and read off the model’s current answer. The probe is a separate forward pass: the controller only *observes*, never halting or altering the main trajectory. This yields, per problem, the full answer trajectory a stopping rule would have access to. This exploratory study uses a deliberately tight 3072 budget so that many trajectories are truncated and the “continue vs. stop” contrast is visible within one A100; the preregistered sweep (Section 3) re-examines the phenomenon at the 16K/32K budgets used for the main results, and we report the recovery statistics there as well.

E.2 Agreement is not correctness

We compare two notions of probe agreement. *Cumulative* agreement (all non-empty probes so far

concur) is nearly trustworthy: on the 87 problems that reach unanimous cumulative agreement, 98.9% are correct, with a single whole-trajectory false consensus. But cumulative agreement is not what an online stopping rule sees. The operative signal is *window* agreement—the last five probes concurring—and it is markedly less reliable: of the 338 problems reaching a unanimous last-five window, only 93.5% are correct, giving a 6.5% (22/338) false-consensus rate. Perfect *global* consensus is almost safe; the *local* consensus a controller must act on is not.

E.3 The cost lands at the stopping decision

We simulate a *naive consecutive-agreement* stop: halt at the first point where three consecutive probes agree, are non-empty, and are *certain* (the model emits the boxed answer with high token-level confidence; formal definition in Appendix A). This is a simplified heuristic, *not* the Certindex algorithm, which we reproduce faithfully in §4.4. It fires on 416/500 problems. The committed answers are correct only 69.2% of the time, whereas letting those same problems run to completion (within the 3072 budget) reaches 85.6%—a 16.4 **pp accuracy loss** in exchange for an average saving of 1321 tokens, with 128 problems stopped on an outright wrong answer.

E.4 Recovery is common, and early stopping kills it

The loss is driven by *recovery*: continued reasoning overturns the intermediate answer far more often than intuition suggests. Of the 145 problems that at some point formed a three-probe consensus *different* from their final answer, 95 (65.5%) ended up correct; and of the 375 problems whose *first* probe was wrong, 76.3% eventually flipped to correct. Neither “first impression” nor “local consensus” is close to terminal.

Counterintuitively, *later* consensus is *less* trustworthy: consensus formed before 512 tokens is 87.4% correct, but consensus that only forms after 2048 tokens is 58.1% correct. Hard problems both converge slowly and converge to worse answers—exactly the problems on which an aggressive stopping rule does most damage.

E.5 Miscalibration

Window share is a poorly calibrated confidence estimate. Its calibration error—the gap between share and empirical accuracy, averaged over share

bins (CCE)—is 0.080 (versus 0.149 for cumulative share), but the error is concentrated where a rule would act: at window share 0.8 the empirical accuracy is only 47.2% ($n=36$). A threshold tuned to “looks nearly unanimous” is operating precisely in the band where the signal is least reliable.

E.6 An error taxonomy

Hand-inspecting stopped-but-wrong cases (28 cases from the first 100 problems; AI-assisted labeling, treated as preliminary) yields *four* non-empty types out of five categories. By count: **(A)** numeric collapse—stable convergence to a wrong number—14/28; **(D)** derivation gaps—dropped roots/cases, unverified solutions—7/28; **(E)** format/option hallucination—a non-multiple-choice problem stably emitting a letter such as “B”—6/28, which is an artifact of the probe mechanism itself rather than a model belief; and **(B)** expression collapse—1/28. Type **(C)** (sign errors) did not occur in this sample. Type E is itself a methodological finding: roughly one in five “false consensus” is a probe formatting artifact, implying that any deployed rule should filter answer *shape*.

Takeaway. These observations—local consensus unreliable, recovery common, late consensus worse, confidence miscalibrated—suggest that a good stopping rule must be much more than “agreement.” The rest of the paper asks whether the *best possible* such rule, drawn from a large preregistered space, can be both safe and economical. In the searched space, it cannot.

F Exploratory: A Boundary-Confidence Controller

Section 4.4 showed that the one prior method staying near full accuracy, DEER, reads a different signal. As a first step beyond the main (negative) result, we push that signal into a deployment-style online controller and find it recovers substantial savings at near-neutral accuracy on dev. We stress up front that these results are *exploratory*—three dev seeds, outside the preregistered sweep, seed-sensitive, and without test-split confirmation—and we report them only to establish that the failure is about the *signal*, not early exit. They are not a validated method and are not part of the paper’s claims.

F.1 From trial-answer confidence to a verification branch

DEER emits a short trial answer at reasoning-boundary (“Wait”) tokens and stops once the model’s confidence in that trial answer (the mean probability of its answer tokens) exceeds a threshold. On our development trajectories the mechanism has a flaw: DEER measures confidence on the trial answer but then *delivers a separately generated readout*, and the two disagree on 14.6% of triggered problems. The extra readout averages 470 tokens yet yields no net accuracy gain (trial 88.9% vs. readout 88.5%), though it corrects some cases on some models. We therefore evaluate an *inspired* controller that (i) *fast-commits* the trial answer when confidence is very high (> 0.995), skipping the redundant readout, and (ii) otherwise opens a *retained verification branch* that commits only if a second high-confidence answer is mathematically equivalent, and otherwise keeps reasoning. Both run in a single online engine that generates the main trajectory live and probes only at scheduled boundaries—no frozen suffixes.

F.2 Online deployment results

Table 7 reports online runs over the 18-environment development set at three seeds (42/43/44; 1,368 trajectories). Across seeds the inspired controller holds macro accuracy near-neutral (-0.75 pp; per-seed range $[-6.1, +4.2]$ pp) while saving 34% of tokens (per-seed 25–44%), still dominating a faithful online DEER reference (-2.7 pp, 22%). A stratified paired bootstrap over the 18 environments (10,000 resamples) gives a token-saving advantage of $+12.1\%$ (95% CI $[+0.7, +22.9]$) at an accuracy difference of $+2.0$ pp (95% CI $[-5.0, +9.0]$): the controller saves significantly more at statistically indistinguishable accuracy. The single-seed 42 point ($+2.8$ pp on Qwen3-8B, 44% saving) sits at the favorable end of this spread.

F.3 Relation to the preregistration

This controller is *not* part of the preregistered consensus sweep: it uses a different signal, a separate configuration, and its own results directory, and it would require its own train/dev $\rightarrow$ test validation before any confirmed claim. It does not modify the sweep’s candidates, gates, or selected rules (there were none). We therefore present it as a positive *signal*, not a finished method: it shows that stepping outside the consensus-signal space escapes

Method	Model	Δ acc	Fair saving
Inspired	Qwen3-8B	-1.5 pp	41.0%
Inspired	DeepSeek-7B	$+0.0$ pp	27.4%
Inspired	macro	-0.75 pp	34.2%
DEER ref.	macro	-2.7 pp	22.1%

Table 7: Online boundary-confidence early exit (dev, seeds 42/43/44, 1,368 trajectories; macro over 18 environments, per-model over three seeds). The inspired fast-path + verification-branch controller keeps macro accuracy near-neutral (-0.75 pp, per-seed range $[-6.1, +4.2]$) while saving 34% of tokens, and saves significantly more than a faithful online DEER reference at indistinguishable accuracy. Exploratory: seed-sensitive, small AMC23/AIME24 sets.

the accuracy tax that Section 4 proved unavoidable inside it.

1030
1031