



Out-of-sample generalization (test split) across scale and architecture: dev-selected DEER points (C/B/T) stay in the gate on every model; consensus train candidates do not

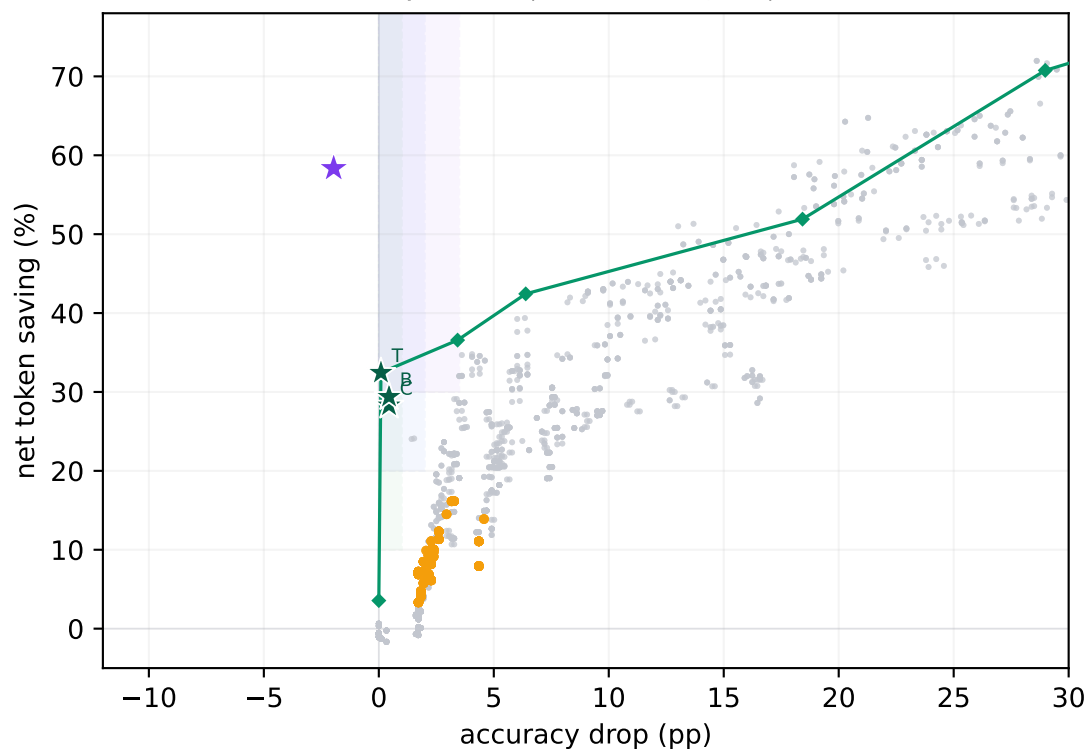
DeepSeek-7B (dev model, test split)



Qwen3-8B (dev model, test split)



Qwen-32B (held-out scale, test)



Llama-8B (held-out arch, test)

